

HOMEWORK 3: 3D RECONSTRUCTION

16-820 Advanced Computer Vision (Fall 2024)

<https://16820advancedcv.github.io/>

OUT: October 3rd, 2024

DUE: October 23rd, 2024

Instructor: Matthew O'Toole

TAs: Nikhil Keetha, Ayush Jain, Yuyao Shi

Instructions/Hints

- Please refer to the [course logistics page](#) for information on the **Collaboration Policy** and **Late Submission Policy**.
- **Submitting your work:** There will be two submission slots for this homework on **Gradescope**: Written and Programming.
 - **Write-up.** For written problems such as short answers, multiple choice, derivations, proofs, or plots, we will be using the written submission slot. Please use this provided template. **We don't accept handwritten submissions.** Each answer should be completed in the boxes provided below the question. You are allowed to adjust the size of these boxes, but **make sure to link your answer to each question when submitting to Gradescope.** Otherwise, your submission will not be graded. To use the provided template - upload the template .zip file directly to [Overleaf](#).
 - **Code.** You are also required to upload your code, which you wrote to solve this homework, to the Programming submission slot. Your code may be run by TAs so please make sure it is in a workable state. The assignment must be completed using Python 3.10.12. We recommend setting up python virtual environment (conda or venv) for the assignment.
 - Regrade requests can be made after the homework grades are released, however, this gives the TA the opportunity to regrade your entire paper, meaning if additional mistakes are found then points will be deducted.
- **Start early!** This homework is difficult and may take a long time to complete.
- **Verify your implementation as you proceed.** If you don't verify that your implementation is correct on toy examples, you will risk having a huge mess when you put everything together.
- **Q&A.** If you have any questions or need clarifications, please post in Slack or visit the TAs during office hours. Additionally, we provide a **FAQ** ([section 8](#)) with questions from previous semesters. Make sure you read it prior to starting your implementations.

Overview

In this assignment, you will be implementing an algorithm to reconstruct a 3D point cloud from a pair of images taken at different angles. In **Part I** you will answer theory questions about 3D reconstruction. In **Part II** you will apply the 8-point algorithm and triangulation to find and visualize 3D locations of corresponding image points.

Part I

Theory

Before implementing our own 3D reconstruction, let's take a look at some simple theory questions that may arise. The answers to the below questions should be relatively short, consisting of a few lines of math and text (maybe a diagram if it helps your understanding).

Q1.1 [5 points] Suppose two cameras fixate on a point \mathbf{x} (see [Figure 1](#)) in space such that their principal axes intersect at that point. Show that if the image coordinates are normalized so that the coordinate origin $(0, 0)$ coincides with the principal point, the F_{33} element of the fundamental matrix is zero.

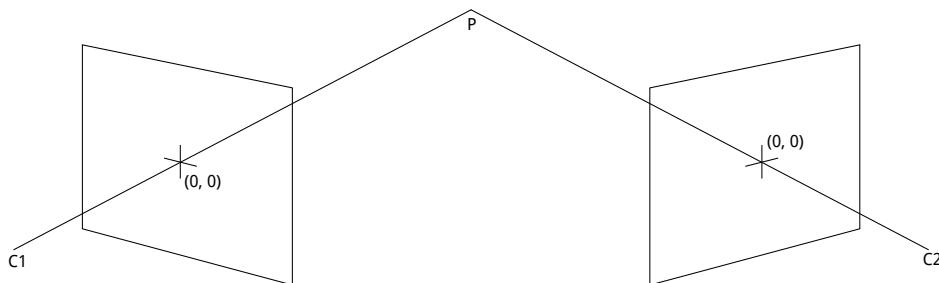


Figure 1: Figure for Q1.1. C_1 and C_2 are the optical centers. The principal axes intersect at point w (P in the figure).

Q1.1

Under normalized image coordinates, both the principal axes of camera C_1 and C_2 intersect at a 3D point w , plus C_1 intercepting at image plane at coordinate $x_1 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$ (in homogeneous coordinate) and C_2 intercepting at image plane at coordinate $x_2 = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$ (in homogeneous coordinate), then x_1 and x_2 satisfy the following equation: $x_2^T F x_1 = 0$, where $F = \begin{bmatrix} F_{11} & F_{12} & F_{13} \\ F_{21} & F_{22} & F_{23} \\ F_{31} & F_{32} & F_{33} \end{bmatrix}$. When we plug in x_1 and x_2 , it shows that $\begin{bmatrix} 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} F_{11} & F_{12} & F_{13} \\ F_{21} & F_{22} & F_{23} \\ F_{31} & F_{32} & F_{33} \end{bmatrix} \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} = 0$, and it results in $F_{33} = 0$.

Q1.2 [5 points] Consider the case of two cameras viewing an object such that the second camera differs from the first by a *pure translation* that is parallel to the x -axis. Show that the epipolar lines in the two cameras are also parallel. Back up your argument with relevant equations. You may assume both cameras have the same intrinsics.

Q1.2

Since the two cameras only differ by a pure translation, the rotation matrix $R = I$ (Identity matrix), and the translation matrix $T = [T_x, 0, 0]^T$. Then we have the following equation: $[x_2 \quad y_2 \quad 1] K^{-\top} E K^{-1} \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} = 0$, where $E = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -T_x \\ 0 & T_x & 0 \end{bmatrix}$. We can view $K^{-1} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x_n \\ y_n \\ 1 \end{bmatrix}$ in normalized camera coordinates, then the above equation would be $[x_{n2} \quad y_{n2} \quad 1] E \begin{bmatrix} x_{n1} \\ y_{n1} \\ 1 \end{bmatrix} = 0$. We now have epipolar line in C_2 as $l_2 = E \begin{bmatrix} x_{n1} \\ y_{n1} \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ -T_x \\ T_x y_{n1} \end{bmatrix} \rightarrow x_{n2}^T \begin{bmatrix} 0 \\ -T_x \\ T_x y_{n1} \end{bmatrix} = -y_{n2} T_x + T_x y_{n1} = 0$, and epipolar line in C_1 as $l_1 = E^T \begin{bmatrix} x_{n2} \\ y_{n2} \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ T_x \\ -T_x y_{n2} \end{bmatrix} \rightarrow x_{n1}^T \begin{bmatrix} 0 \\ T_x \\ -T_x y_{n2} \end{bmatrix} = y_{n1} T_x - T_x y_{n2} = 0$. We can see that l_1 and l_2 are both horizontal lines (parallel to the x -axis), and they are parallel to each other.

Q1.3 [5 points] Suppose we have an inertial sensor that gives us the accurate positions (\mathbf{R}_i and \mathbf{t}_i , the rotation matrix and translation vector) of the robot at time i . What will be the effective rotation (\mathbf{R}_{rel}) and translation (\mathbf{t}_{rel}) between two frames at different time stamps? Suppose the camera intrinsics (\mathbf{K}) are known, express the essential matrix (\mathbf{E}) and the fundamental matrix (\mathbf{F}) in terms of \mathbf{K} , \mathbf{R}_{rel} and \mathbf{t}_{rel} .

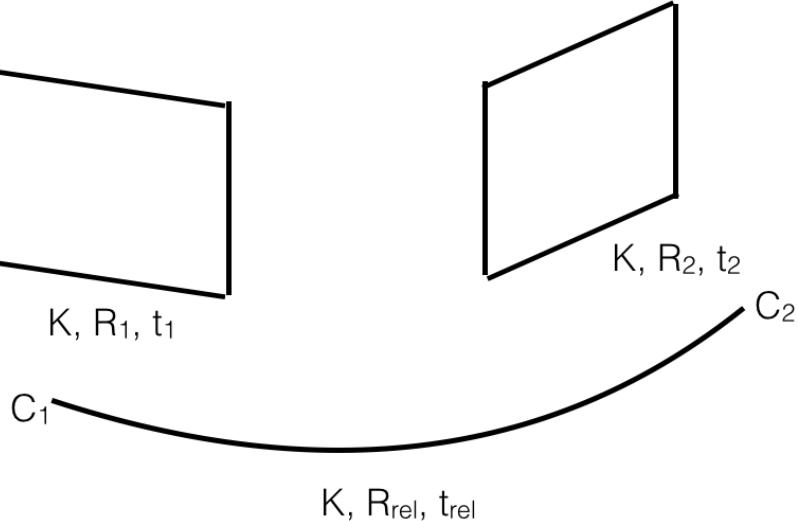


Figure 2: Figure for Q1.3. C_1 and C_2 are the optical centers. The rotation and the translation is obtained using inertial sensors. \mathbf{R}_{rel} and \mathbf{t}_{rel} are the relative rotation and translation between two frames.

Q1.3

Assume $P = \begin{bmatrix} X_P \\ Y_P \\ Z_P \end{bmatrix}$ is the 3D interception point of camera at time 1 and time 2, that is, camera C1 and C2. Then we have image point in C1 and C2 as x_1 and x_2 as the following equations:

$$\lambda_1 \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} = K(R_1 \begin{bmatrix} X_P \\ Y_P \\ Z_P \end{bmatrix}) + t_1 \quad (1)$$

$$\lambda_2 \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} = K(R_2 \begin{bmatrix} X_P \\ Y_P \\ Z_P \end{bmatrix}) + t_2 \quad (2)$$

The λ_1 and λ_2 are the scalar for homogeneous coordinates. When we substitute $\begin{bmatrix} X_P \\ Y_P \\ Z_P \end{bmatrix}$ with $\begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix}$ in equation (2), then we have the following equation:

$$\begin{aligned} \lambda_2 \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix} &= K(R_2 R_1^{-1}(K^{-1} \lambda_1 \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} - t_1) + t_2) \\ &= KR_2 R_1^{-1} K^{-1} \lambda_1 \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} - KR_2 R_1^{-1} t_1 + Kt_2 \\ &= KR_{rel} K^{-1} \lambda_1 \begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} + Kt_{rel} \end{aligned} \quad (3)$$

Accordingly, $R_{rel} = R_2 R_1^{-1}$, and $t_{rel} = t_2 - R_2 R_1^{-1} t_1$. Also, the essential matrix E is the cross product of t_{rel} and R_{rel} : $E = t_{rel} \times R_{rel}$, and the fundamental matrix F is: $F = K^{-T} (t_{rel} \times R_{rel}) K^{-1}$.

Part II

Practice

1 Overview

In this part you will begin by implementing the 8-point algorithm seen in class to estimate the fundamental matrix from corresponding points in two images ([section 2](#)). Next, given the fundamental matrix and calibrated intrinsics (which will be provided) you will compute the essential matrix and use this to compute a 3D metric reconstruction from 2D correspondences using triangulation ([section 3](#)). Then, you will implement a method to automatically match points taking advantage of epipolar constraints and make a 3D visualization of the results ([section 4](#)). Finally, you will implement RANSAC and bundle adjustment to further improve your algorithm ([section 5](#)).

2 Fundamental Matrix Estimation

In this section you will explore different methods of estimating the fundamental matrix given a pair of images. In the `data/` directory, you will find two images (see [Figure 3](#)) from the Middlebury multi-view dataset¹, which is used to evaluate the performance of modern 3D reconstruction algorithms.

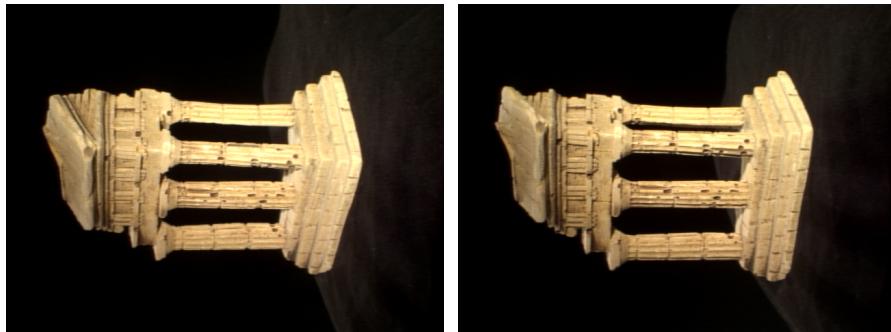


Figure 3: Temple images for this assignment

2.1 The Eight Point Algorithm

The 8-point algorithm (discussed in class, and outlined in Section 8.1 of [[1](#)]) is arguably the simplest method for estimating the fundamental matrix. For this section, you can use provided correspondences you can find in `data/some_corresp.npz`.

Q2.1 [10 points] Finish the function `eightpoint` in `q2_1_eightpoint.py`. Make sure you follow the signature for this portion of the assignment:

```
F = eightpoint(pts1, pts2, M)
```

where `pts1` and `pts2` are $N \times 2$ matrices corresponding to the (x, y) coordinates of the N points in the first and second image respectively. `M` is a scale parameter.

¹<http://vision.middlebury.edu/mview/data/>

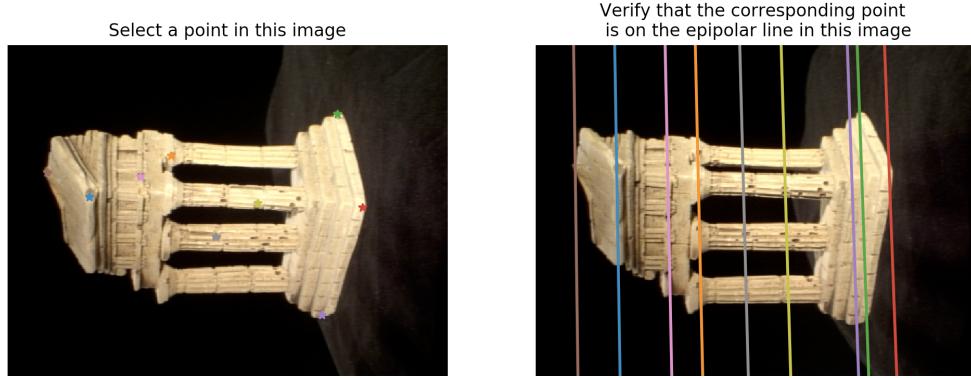


Figure 4: `displayEpipolarF` in `helper.py` creates a GUI for visualizing epipolar lines

- You should scale the data as was discussed in class, by dividing each coordinate by M (the maximum of the image's width and height). After computing \mathbf{F} , you will have to “unscale” the fundamental matrix.

Hint: If $\mathbf{x}_{normalized} = \mathbf{T}\mathbf{x}$, then $\mathbf{F}_{unnormalized} = \mathbf{T}^T\mathbf{FT}$.

You must enforce the singularity condition of \mathbf{F} before unscaling.

- You may find it helpful to refine the solution by using local minimization. This probably won't fix a completely broken solution, but may make a good solution better by locally minimizing a geometric cost function. For this we have provided a helper function `refineF` in `helper.py` taking in \mathbf{F} and the two sets of points, which you can call from `eightpoint` before unscaling \mathbf{F} .
- Remember that the x -coordinate of a point in the image is its column entry, and y -coordinate is the row entry. Also note that eight-point is just a figurative name, it just means that you need at least 8 points; your algorithm should use an over-determined system ($N > 8$ points).
- To visualize the correctness of your estimated \mathbf{F} , use the function `displayEpipolarF` in `helper.py`, which takes in \mathbf{F} , and the two images. This GUI lets you select a point in one of the images and visualize the corresponding epipolar line in the other image (Figure 4).
- In addition to visualization, we also provide a test code snippet in `q2_1_eightpoint.py` which uses helper function `calc_epi_error` to evaluate the quality of the estimated fundamental matrix. This function calculates the distance between the estimated epipolar line and the corresponding points. For the eight point algorithm, the error should on average be < 1 .

Output: Save your matrix \mathbf{F} and scale \mathbf{M} to the file `q2_1.npz`.

In your write-up:

- Write your recovered \mathbf{F}
- Include an image of some example output of `displayEpipolarF`
- Include the code snippet of `eightpoint` function

Q2.1

As in the output: `./q2_1.npz`, the recovered F is:

$$\begin{pmatrix} -2.19293792e^{-07} & 2.95926413e^{-05} & -2.51886251e^{-01} \\ 1.28064423e^{-05} & -6.64493522e^{-07} & 2.63771761e^{-03} \\ 2.42228993e^{-01} & -6.82585388e^{-03} & 1.00000000e^{+00} \end{pmatrix}$$

Besides, M = 640.

The following [Figure 5](#) shows the result of the Eight Point Algorithm:

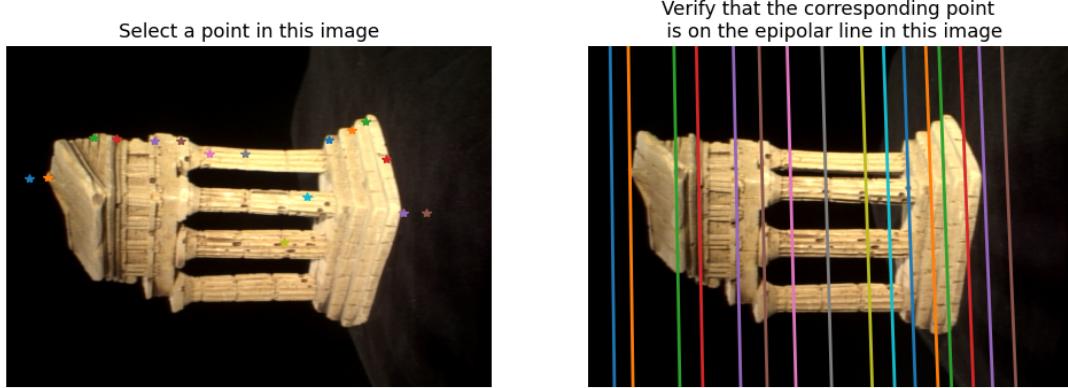


Figure 5: Result of the Eight Point Algorithm

The following [Figure 6](#) and [Figure 7](#) show the code snippet of the Eight Point Algorithm in `q2_1_eightpoint.py`:

```

67 """
68 Q2.1: Eight Point Algorithm
69 Input: pts1, Nx2 Matrix
70     pts2, Nx2 Matrix
71     M, a scalar parameter computed as max (imwidth, imheight)
72 Output: F, the fundamental matrix
73
74 HINTS:
75 (1) Normalize the input pts1 and pts2 using the matrix T.
76 (2) Setup the eight point algorithm's equation.
77 (3) Solve for the least square solution using SVD.
78 (4) Use the function 'singularize' (provided) to enforce the singularity condition.
79 (5) Use the function 'refineF' (provided) to refine the computed fundamental matrix.
80     (Remember to use the normalized points instead of the original points)
81 (6) Unscale the fundamental matrix
82 """
83
84 def eightpoint(pts1, pts2):
85     # Replace pass by your implementation
86     # -----
87     # YOUR CODE HERE
88
89     # Normalizing to [-1, 1] x [-1, 1]
90     T_matrix, pts1_norm, pts2_norm = normalizeImagePts(pts1, pts2, M)
91
92     # Setup equations and solve least square solution using SVD
93     F_matrix = computeF(pts1_norm, pts2_norm)
94
95     # Enforce singularity condition.
96     F_sing = _singularize(F_matrix)
97
98     # Refined the F
99     F_refine = refineF(F_sing, pts1_norm, pts2_norm)
100    #F_refine = F_sing
101
102    # Unscale the F
103    F_unscaled = T_matrix.T @ F_refine @ T_matrix
104
105    # Make sure F is unique avoiding infinite scaling possibilities
106    F_ret = F_unscaled / F_unscaled[2, 2]
107
108    return F_ret

```

Figure 6: First Code Snippet

```

8     def computeF(x1, x2):
9         # Construct A matrix
10        A = np.zeros((x1.shape[0], 9))
11        for i in range(x1.shape[0]):
12            x11 = x1[i][0]
13            y11 = x1[i][1]
14            x12 = x1[i][0]
15            y12 = x1[i][1]
16            A[i] = [x11*x12, x12*y11, x12, y12*x11, y12*y11, y12, x11, y11, 1]
17
18        # Solve by directly call SVD
19        try:
20            u, sigma, v = np.linalg.svd(A)
21        except np.linalg.LinAlgError as e:
22            eps = 1e-10 # Small regularization constant
23            perturbation = eps * np.random.rand(A.shape[0], A.shape[1])
24            A_regularized = A + perturbation
25            u, sigma, v = np.linalg.svd(A_regularized)
26
27        # Choose the eigenvector w/ smallest eigenvalues
28        v1_smallest_eigenvec = v[-1, :]
29        F2to1_v1_smallest = np.reshape(v1_smallest_eigenvec, (3, 3))
30        F2to1_v1 = np.array(F2to1_v1_smallest, dtype=np.float32)
31
32        return F2to1_v1
33
34    def normalizeImagePts(pts1, pts2, M):
35        T_matrix = np.zeros((3, 3), dtype=np.float32)
36        T_matrix[0][0] = 1/M
37        T_matrix[0][2] = 0
38        T_matrix[1][1] = 1/M
39        T_matrix[1][2] = 0
40        T_matrix[2][2] = 1
41
42        ...
43        T_matrix = np.zeros((3, 3), dtype=np.float32)
44        T_matrix[0][0] = 2/M
45        T_matrix[0][2] = -1
46        T_matrix[1][1] = 2/M
47        T_matrix[1][2] = -1
48        T_matrix[2][2] = 1
49
50
51        x1_homo = np.hstack([pts1, np.ones((pts1.shape[0], 1))]) #Nx3
52        x1_norm = T_matrix @ x1_homo.T #Nx3
53        x1_norm = x1_norm.T #Nx3
54        x1_col1 = x1_norm[:, 0]/x1_norm[:, 2]
55        x1_col2 = x1_norm[:, 1]/x1_norm[:, 2]
56        x1_norm = np.column_stack((x1_col1, x1_col2)) #Nx2
57
58        x2_homo = np.hstack([pts2, np.ones((pts2.shape[0], 1))]) #Nx3
59        x2_norm = T_matrix @ x2_homo.T #Nx3
60        x2_norm = x2_norm.T #Nx3
61        x2_col1 = x2_norm[:, 0]/x2_norm[:, 2]
62        x2_col2 = x2_norm[:, 1]/x2_norm[:, 2]
63        x2_norm = np.column_stack((x2_col1, x2_col2)) #Nx2
64
65        return T_matrix, x1_norm, x2_norm

```

Figure 7: Second Code Snippet

2.2 The Seven Point Algorithm (Extra Credit)

Since the fundamental matrix only has seven degrees of freedom, it is possible to calculate \mathbf{F} using only seven-point correspondences. This requires solving a polynomial equation. In this section, you will implement the seven-point algorithm (outlined in this [post](#)).

Q2.2 [Extra Credit - 15 points] Finish the function `sevenpoint` in `q2_2_sevenpoint.py`. Make sure you follow the signature for this portion of the assignment:

```
Farray = sevenpoint(pts1, pts2, M)
```

where `pts1` and `pts2` are 7×2 matrices containing the correspondences and `M` is the normalizer (use the maximum of the image's height and width), and `Farray` is a list array of length either 1 or 3 containing Fundamental matrix/matrices. Use `M` to normalize the point values between $[0, 1]$ and remember to “unnormalize” your computed \mathbf{F} afterward.

Manually select 7 points from the provided point in `data/some_corresp.npz`, and use these points to recover a fundamental matrix \mathbf{F} . Use `calc_epi_error` in `helper.py` to calculate the error to pick the best one, and use `displayEpipolarF` to visualize and verify the solution.

Output: Save your matrix \mathbf{F} and scale \mathbf{M} to the file `q2_2.npz`.

In your write-up:

- Write your recovered \mathbf{F}
- Include an image of some example output of `displayEpipolarF`
- Include the code snippet of `sevenpoint` function

Q2.2

As in the output: `./q2_2.npz`, the recovered F is:

$$\begin{pmatrix} 8.10457899e-07 & 8.90918106e-06 & -2.01028454e-01 \\ 2.63329911e-05 & -6.00542406e-07 & 6.97427365e-04 \\ 1.92182078e-01 & -4.20123348e-03 & 1.00000000e+00 \end{pmatrix}$$

Besides, $M = 640$.

The following [Figure 8](#) shows the result of the Eight Point Algorithm:

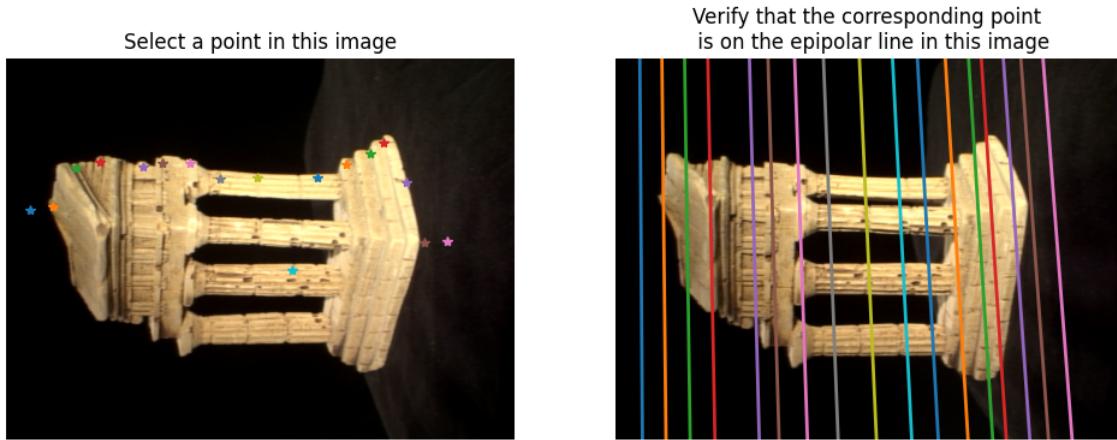


Figure 8: Result of the Seven Point Algorithm

The following [Figure 9](#), [Figure 10](#), [Figure 11](#) and [Figure 12](#) show the code snippet of the Seven Point Algorithm in `q2_2_sevenpoint.py`:

```

100 """
101 Q2.2: Seven Point Algorithm for calculating the fundamental matrix
102     Input: pts1, 2x2 Matrix containing the corresponding points from image1
103         pts2, 2x2 Matrix containing the corresponding points from image2
104     M, a scalar parameter computed as max (imwidth, imheight)
105     Output: Farray, a list of estimated 3x3 fundamental matrices.
106
107 HINTS:
108 (1) Normalize the input pts1 and pts2 scale parameter M.
109 (2) Use the 8-point algorithm.
110 (3) Compute the least square solution using SVD.
111 (4) Pick the last two column vector of v1.T (the two null space solution f1 and f2)
112 (5) Use the singularity constraint to solve for the cubic polynomial equation of F = a*f1 + (1-a)*f2 that leads to
113     det(F) = 0. Solving this polynomial will give you one or three real solutions of the fundamental matrix.
114 (6) Use np.polynomial.polynomial.polyroots to solve for the roots
115 (7) Unscaled the fundamental matrices and return as Farray
116
117 """
118
119 def sevenpoint(pts1, pts2, M):
120     Farray = []
121     # ----- TODO -----
122     # YOUR CODE HERE
123
124     # Normalizing to [0, 1] x [0, 1]
125     T_matrix, pts1_norm, pts2_norm = normalizeImagePts(pts1, pts2, M)
126
127     # Setup matrix and solve SVD and pick the last two columns
128     F_v1_matrix, F_v2_matrix = computeF(pts1_norm, pts2_norm)
129
130     # Calculate cd=0
131     coefficients = calcCoeff(F_v1_matrix, F_v2_matrix)
132
133     # Solve for the roots
134     roots = np.polynomial.polynomial.polyroots(list(coefficients))
135
136     # Set up equations and solve least square solution using SVD
137     F_alpha_array = [singularize(alpha, F_v1_matrix, F_v2_matrix) for alpha in roots]
138
139     # Enforce singularity condition.
140     F_sing_array = []
141     for F_matrix in F_alpha_array:
142         try:
143             F_sing = singularize(F_matrix)
144             except np.linalg.LinAlgError as e:
145                 F_sing = F_matrix.copy()
146                 F_sing[0][0] = F_sing[1][1]
147                 F_sing[1][1] = F_sing[0][0]
148             # F_sing_array = [singularize(F_matrix) for F_matrix in F_alpha_array]
149             F_sing_array = [F_matrix for F_matrix in F_alpha_array]
150
151     # Refine the F
152     F_refine_array = []
153     for F_sing in F_sing_array:
154         try:
155             F_refine = refine(F_sing, pts1_norm, pts2_norm) for F_sing in F_sing_array
156             except np.linalg.LinAlgError as e:
157                 F_refine = F_sing.copy()
158             F_refine_array.append(F_refine)
159             # F_refine_array = [refine(F_sing, pts1_norm, pts2_norm) for F_sing in F_sing_array]
160             F_refine_array = [F_sing for F_sing in F_sing_array]
161
162     # Unsacle the F
163     F_unscaled_array = [T_matrix.T @ F_refine @ T_matrix for F_refine in F_refine_array]
164
165     # Make sure F is unique avoiding infinite scaling possibilities.
166     Farray = [F_unscaled/F_unscaled[0, 0] for F_unscaled in F_unscaled_array]
167
168 return Farray

```

Figure 10: Second Code Snippet

Figure 9: First Code Snippet

Q2.2 continued

```

32 def computeF(x1, x2):
33     # Construct A matrix
34     A = np.zeros((x1.shape[0], 9))
35     for i in range(x1.shape[0]):
36         x1i = x1[i][0]
37         y1i = x1[i][1]
38         x2i = x2[i][0]
39         y2i = x2[i][1]
40         A[i] = [x1i*x12, x12*y1i, x12, yi2*x1i, yi1*y1i, yi2, x1i, yi1, 1]
41
42     # Solve by directly call SVD
43     try:
44         u, sigma, v = np.linalg.svd(A)
45     except np.linalg.LinAlgError as e:
46         eps = 1e-10 # Small regularization constant
47         perturbation = eps * np.random.rand(A.shape[0], A.shape[1])
48         A_regularized = A + perturbation
49         u, sigma, v = np.linalg.svd(A_regularized)
50
51     v1_smallest_eigenvec = v[:, -1, :]
52     v2_smallest_eigenvec = v[:, -2, :]
53
54     F2to1_v1_smallest = np.reshape(v1_smallest_eigenvec, (3, 3))
55     F2to1_v2_smallest = np.reshape(v2_smallest_eigenvec, (3, 3))
56     F2to1_v1 = np.array(F2to1_v1_smallest, dtype=np.float32)
57     F2to1_v2 = np.array(F2to1_v2_smallest, dtype=np.float32)
58
59     return F2to1_v1, F2to1_v2

```

Figure 11: Third Code Snippet

```

61 def normalizeImagePts(pts1, pts2, M):
62     T_matrix = np.zeros((3, 3), dtype=np.float32)
63     T_matrix[0][0] = 1/M
64     T_matrix[0][2] = 0
65     T_matrix[1][1] = 1/M
66     T_matrix[1][2] = 0
67     T_matrix[2][2] = 1
68
69     ...
70
71     T_matrix[0][0] = 2/M
72     T_matrix[0][2] = -1
73     T_matrix[1][1] = 2/M
74     T_matrix[1][2] = -1
75     T_matrix[2][2] = 1
76
77
78     x1_homo = np.hstack([pts1, np.ones((pts1.shape[0], 1))]) #Nx3
79     x1_norm = T_matrix @ x1_homo.T #3xN
80     x1_norm = x1_norm.T #Nx3
81     x1_col1 = x1_norm[:, 0]/x1_norm[:, 2]
82     x1_col2 = x1_norm[:, 1]/x1_norm[:, 2]
83     x1_norm = np.column_stack((x1_col1, x1_col2)) #Nx2
84
85     x2_homo = np.hstack([pts2, np.ones((pts2.shape[0], 1))]) #Nx3
86     x2_norm = T_matrix @ x2_homo.T #3xN
87     x2_norm = x2_norm.T #Nx3
88     x2_col1 = x2_norm[:, 0]/x2_norm[:, 2]
89     x2_col2 = x2_norm[:, 1]/x2_norm[:, 2]
90     x2_norm = np.column_stack((x2_col1, x2_col2)) #Nx2
91
92
93     return T_matrix, x1_norm, x2_norm

```

Figure 12: Fourth Code Snippet

3 Metric Reconstruction

You will compute the camera matrices and triangulate the 2D points to obtain the 3D scene structure. To obtain the Euclidean scene structure, first convert the fundamental matrix \mathbf{F} to an essential matrix \mathbf{E} . Examine the lecture notes and the textbook to find out how to do this when the internal camera calibration matrices \mathbf{K}_1 and \mathbf{K}_2 are known; these are provided in `data/intrinsics.npz`.

Q3.1 [5 points] Complete the function `essentialMatrix` in `q3_1_essential_matrix.py` to compute the essential matrix \mathbf{E} given \mathbf{F} , \mathbf{K}_1 and \mathbf{K}_2 with the signature:

```
E = essentialMatrix(F, K1, K2)
```

Output: Save your estimated \mathbf{E} using \mathbf{F} from the eight-point algorithm to `q3_1.npz`.
In your write-up:

- Write your estimated \mathbf{E}
- Include the code snippet of `essentialMatrix` function

Q3.1

As in the output: `./q3_1.npz`, the recovered \mathbf{E} is:

$$\begin{pmatrix} -5.06923074e^{-01} & 6.86542875e^{+01} & -3.71961318e^{+02} \\ 2.97106690e^{+01} & -1.54718732e^{+00} & 9.68232032e^{+00} \\ 3.72990948e^{+02} & 2.98549953e^{+00} & 1.50354471e^{-01} \end{pmatrix}$$

The following **Figure 13** shows the code snippet of the `q3_1_essential_matrix.py`:

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 from q2_1_eightpoint import eightpoint
5
6 # Insert your package here
7
8 """
9 Q3.1: Compute the essential matrix E.
10 Input: F, fundamental matrix
11          K1, internal camera calibration matrix of camera 1
12          K2, internal camera calibration matrix of camera 2
13 Output: E, the essential matrix
14 """
15
16
17 def essentialMatrix(F, K1, K2):
18     # Replace pass by your implementation
19     # ----- TODO -----
20     # YOUR CODE HERE
21     return K2.T @ F @ K1
22
23
24 if __name__ == "__main__":
25     correspondence = np.load("data/some_corresp.npz") # Loading correspondences
26     intrinsics = np.load("data/intrinsics.npz") # Loading the intrinsics of the camera
27     K1, K2 = intrinsics["K1"], intrinsics["K2"]
28     pts1, pts2 = correspondence["pts1"], correspondence["pts2"]
29     im1 = plt.imread("data/im1.png")
30     im2 = plt.imread("data/im2.png")
31
32     F = eightpoint(pts1, pts2, M=np.max([*im1.shape, *im2.shape]))
33     E = essentialMatrix(F, K1, K2)
34     np.savez("q3_1.npz", E, F)
35
36
37 # Simple Tests to verify your implementation:
38 assert np.linalg.matrix_rank(E) == 2
39

```

Figure 13: Code Snippet

Given an essential matrix, it is possible to retrieve the projective camera matrices \mathbf{M}_1 and \mathbf{M}_2 from it. Assuming \mathbf{M}_1 is fixed at $[\mathbf{I}, 0]$, \mathbf{M}_2 can be retrieved up to a scale and four-fold rotation ambiguity. For

details on recovering \mathbf{M}_2 , see section 11.3 in Szeliski. We have provided you with the function `camera2` in `python/helper.py` to recover the four possible \mathbf{M}_2 matrices given \mathbf{E} .

Note: The matrices \mathbf{M}_1 and \mathbf{M}_2 here are of the form: $\mathbf{M}_1 = [\mathbf{I}|0]$ and $\mathbf{M}_2 = [\mathbf{R}|\mathbf{t}]$.

Q3.2 [10 points] Using the above, complete the function `triangulate` in `q3_2_triangulate.py` to triangulate a set of 2D coordinates in the image to a set of 3D points with the signature:

```
[w, err] = triangulate(C1, pts1, C2, pts2)
```

where pts1 and pts2 are the $N \times 2$ matrices with the 2D image coordinates and w is an $N \times 3$ matrix with the corresponding 3D points per row. $C1$ and $C2$ are the 3×4 camera matrices. Remember that you will need to multiply the given intrinsics matrices with your solution for the canonical camera matrices to obtain the final camera matrices. Various methods exist for triangulation - probably the most familiar for you is based on least squares (see [2] Chapter 7 if you want to learn about other methods).

For each point i , we want to solve for 3D coordinates $\mathbf{w}_i = [x_i, y_i, z_i]^T$, such that when they are projected back to the two images, they are close to the original 2D points. To project the 3D coordinates back to 2D images, we first write \mathbf{w}_i in homogeneous coordinates, and compute $\mathbf{C}_1\tilde{\mathbf{w}}_i$ and $\mathbf{C}_2\tilde{\mathbf{w}}_i$ to obtain the 2D homogeneous coordinates projected to camera 1 and camera 2, respectively.

For each point i , we can write this problem in the following form:

$$\mathbf{A}_i \mathbf{w}_i = 0,$$

where \mathbf{A}_i is a 4×4 matrix, and $\tilde{\mathbf{w}}_i$ is a 4×1 vector of the 3D coordinates in the homogeneous form. Then, you can obtain the homogeneous least-squares solution (discussed in class) to solve for each \mathbf{w}_i .

Once you have implemented triangulation, check the performance by looking at the reprojection error:

$$\text{err} = \sum_i \|\mathbf{x}_{1i}, \hat{\mathbf{x}}_{1i}\|^2 + \|\mathbf{x}_{2i}, \hat{\mathbf{x}}_{2i}\|^2$$

where $\hat{\mathbf{x}}_{1i} = \text{Proj}(\mathbf{C}_1, \mathbf{w}_i)$ and $\hat{\mathbf{x}}_{2i} = \text{Proj}(\mathbf{C}_2, \mathbf{w}_i)$. You should see an error less than 500. Ours is around 350.

Note: $C1$ and $C2$ here are projection matrices of the form: $\mathbf{C}_1 = \mathbf{K}_1\mathbf{M}_1 = \mathbf{K}_1[\mathbf{I}|0]$ and $\mathbf{C}_2 = \mathbf{K}_2\mathbf{M}_2 = \mathbf{K}_2[\mathbf{R}|\mathbf{t}]$.

In your write-up:

- Write down the expression for the matrix \mathbf{A}_i for triangulating a pair of 2D coordinates in the image to a 3D point.
- Include the code snippet of `triangulate` function.

Q3.2

Assume that $\mathbf{w}_i = \begin{bmatrix} x_i \\ y_i \\ z_i \end{bmatrix}$ is the 3D coordinates at point i, and $\tilde{\mathbf{w}}_i = \begin{bmatrix} x_i \\ y_i \\ z_i \\ 1 \end{bmatrix}$ is its homogeneous

coordinate. Besides, $u_1 = \lambda_1 \begin{bmatrix} u_{xi_1} \\ u_{yi_1} \\ 1 \end{bmatrix}$ is the 2D projected homogeneous coordinates on camera1 and

$u_2 = \lambda_2 \begin{bmatrix} u_{xi_2} \\ u_{yi_2} \\ 1 \end{bmatrix}$ is the 2D projected homogeneous coordinates on camera2. Let us define the C1 and C2 camera matrix as the following:

$$C1 = \begin{bmatrix} c_{111} & c_{112} & c_{113} & c_{114} \\ c_{121} & c_{122} & c_{123} & c_{124} \\ c_{131} & c_{132} & c_{133} & c_{134} \end{bmatrix} \quad (4)$$

$$C2 = \begin{bmatrix} c_{211} & c_{212} & c_{213} & c_{214} \\ c_{221} & c_{222} & c_{223} & c_{224} \\ c_{231} & c_{232} & c_{233} & c_{234} \end{bmatrix} \quad (5)$$

Then, we have the following two equations:

$$\lambda_1 \begin{bmatrix} u_{xi_1} \\ u_{yi_1} \\ 1 \end{bmatrix} = \begin{bmatrix} c_{111} & c_{112} & c_{113} & c_{114} \\ c_{121} & c_{122} & c_{123} & c_{124} \\ c_{131} & c_{132} & c_{133} & c_{134} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ z_i \\ 1 \end{bmatrix} \quad (6)$$

$$\lambda_2 \begin{bmatrix} u_{xi_2} \\ u_{yi_2} \\ 1 \end{bmatrix} = \begin{bmatrix} c_{211} & c_{212} & c_{213} & c_{214} \\ c_{221} & c_{222} & c_{223} & c_{224} \\ c_{231} & c_{232} & c_{233} & c_{234} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ z_i \\ 1 \end{bmatrix} \quad (7)$$

After multiplying both equation (6) and (7) out, dividing x and y coordinates by the z coordinate, and rearranging the equations, we get the following equation:

$$\begin{bmatrix} c_{111} - u_{xi_1}c_{131} & c_{112} - u_{xi_1}c_{132} & c_{113} - u_{xi_1}c_{133} & c_{114} - u_{xi_1}c_{134} \\ c_{121} - u_{yi_1}c_{131} & c_{122} - u_{yi_1}c_{132} & c_{123} - u_{yi_1}c_{133} & c_{124} - u_{yi_1}c_{134} \\ c_{211} - u_{xi_2}c_{231} & c_{212} - u_{xi_2}c_{232} & c_{213} - u_{xi_2}c_{233} & c_{214} - u_{xi_2}c_{234} \\ c_{221} - u_{yi_2}c_{231} & c_{222} - u_{yi_2}c_{232} & c_{223} - u_{yi_2}c_{233} & c_{224} - u_{yi_2}c_{234} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ z_i \\ 1 \end{bmatrix} = 0 \quad (8)$$

Accordingly, we have A matrix:

$$A = \begin{bmatrix} c_{111} - u_{xi_1}c_{131} & c_{112} - u_{xi_1}c_{132} & c_{113} - u_{xi_1}c_{133} & c_{114} - u_{xi_1}c_{134} \\ c_{121} - u_{yi_1}c_{131} & c_{122} - u_{yi_1}c_{132} & c_{123} - u_{yi_1}c_{133} & c_{124} - u_{yi_1}c_{134} \\ c_{211} - u_{xi_2}c_{231} & c_{212} - u_{xi_2}c_{232} & c_{213} - u_{xi_2}c_{233} & c_{214} - u_{xi_2}c_{234} \\ c_{221} - u_{yi_2}c_{231} & c_{222} - u_{yi_2}c_{232} & c_{223} - u_{yi_2}c_{233} & c_{224} - u_{yi_2}c_{234} \end{bmatrix} \quad (9)$$

Q3.2 continued

The following **Figure 14** shows the code snippet of triangulate() in the q3_2triangulate.py:

```

11 """
12 Q3.2: Triangulate a set of 2D coordinates in the image to a set of 3D points.
13     Input: C1, the 3x4 camera matrix
14             pts1, the Nx2 matrix with the 2D image coordinates per row
15             C2, the 3x4 camera matrix
16             pts2, the Nx2 matrix with the 2D image coordinates per row
17     Output: P, the Nx3 matrix with the corresponding 3D points per row
18             err, the reprojection error.
19
20     Hints:
21     (1) For every input point, form A using the corresponding points from pts1 & pts2 and C1 & C2
22     (2) Solve for the least square solution using np.linalg.svd
23     (3) Calculate the reprojection error using the calculated 3D points and C1 & C2 (do not forget to convert from
24         homogeneous coordinates to non-homogeneous ones)
25     (4) Keep track of the 3D points and projection error, and continue to next point
26     (5) You do not need to follow the exact procedure above.
27 """
28
29
30 def triangulate(C1, pts1, C2, pts2):
31     # Replace pass by your implementation
32     # ----- TODO -----
33     # YOUR CODE HERE
34     # Form Matrix A
35     P = np.zeros((pts1.shape[0], 3), dtype=np.float32) # the reconstructed Nx3 3D points
36     err = 0.0
37     for i in range(pts1.shape[0]):
38         A = np.zeros((4, 4), dtype=np.float32)
39         ux1i = pts1[i, 0]
40         uy1i = pts1[i, 1]
41         ux2i = pts2[i, 0]
42         uy2i = pts2[i, 1]
43
44         A[0] = [C1[0][0]-ux1i*C1[2][0], C1[0][1]-ux1i*C1[2][1], C1[0][2]-ux1i*C1[2][2], C1[0][3]-ux1i*C1[2][3]]
45         A[1] = [C1[1][0]-uy1i*C1[2][0], C1[1][1]-uy1i*C1[2][1], C1[1][2]-uy1i*C1[2][2], C1[1][3]-uy1i*C1[2][3]]
46         A[2] = [C2[0][0]-ux2i*C2[2][0], C2[0][1]-ux2i*C2[2][1], C2[0][2]-ux2i*C2[2][2], C2[0][3]-ux2i*C2[2][3]]
47         A[3] = [C2[1][0]-uy2i*C2[2][0], C2[1][1]-uy2i*C2[2][1], C2[1][2]-uy2i*C2[2][2], C2[1][3]-uy2i*C2[2][3]]
48
49         # Solve by directly call SVD
50         try:
51             u, sigma, v = np.linalg.svd(A)
52         except np.linalg.LinAlgError as e:
53             eps = 1e-10 # Small regularization constant
54             perturbation = eps * np.random.rand(A.shape[0], A.shape[1])
55             A_regularized = A + perturbation
56             u, sigma, v = np.linalg.svd(A_regularized)
57
58         v1_smallest_eigenvect = v[-1, :].reshape(1, -1) # the solution of wi
59         P[i] = v1_smallest_eigenvect[:, 0:3] / v1_smallest_eigenvect[:, -1] # back from homogeneous coordinates to normal coordinates
60
61     # Calculate the error
62     w1_homo = np.hstack([P, np.ones((P.shape[0], 1))]) #Nx4
63     w1_proj = C1 @ w1_homo.T #Nx4x4xN
64     w1_proj = w1_proj.T #Nx3
65     x1_col1 = w1_proj[:, 0] @ w1_proj[:, 2]
66     x1_col2 = w1_proj[:, 1] @ w1_proj[:, 2]
67     x1 = np.column_stack((x1_col1, x1_col2)) #Nx2
68     w2_proj = C2 @ w1_homo.T #Nx4x4xN
69     w2_proj = w2_proj.T #Nx3
70     x2_col1 = w2_proj[:, 0] @ w2_proj[:, 2]
71     x2_col2 = w2_proj[:, 1] @ w2_proj[:, 2]
72     x2 = np.column_stack((x2_col1, x2_col2)) #Nx2
73     err = np.sum(np.linalg.norm(pts1-x1, axis=1)**2) + np.sum(np.linalg.norm(pts2-x2, axis=1)**2)
74
75     return P, err
76

```

Figure 14: Code Snippet

Q3.3 [10 points] Complete the function `findM2` in `q3_2_triangulate.py` to obtain the correct `M2` from `M2s` by testing the four solutions through triangulations.

Use the correspondences from `data/some_corresp.npz`.

Output: Save the correct `M2`, the corresponding `C2`, and 3D points `P` to `q3_3.npz`.

In your writeup: Include the code snippet of `findM2` function.

Q3.3

The following **Figure 15** shows the code snippet of findM2() in the q3_2triangulate.py:

```

86 def findM2(F, pts1, pts2, intrinsics, filename="q3_3.npz"):
87     """
88     Q2.2: Function to find camera2's projective matrix given correspondences
89     Input: F, the pre-computed fundamental matrix
90             pts1, the Nx2 matrix with the 2D image coordinates per row
91             pts2, the Nx2 matrix with the 2D image coordinates per row
92             intrinsics, the intrinsics of the cameras, load from the .npz file
93             filename, the filename to store results
94     Output: [M2, C2, P] the computed M2 (3x4) camera projective matrix, C2 (3x4) K2 * M2, and the 3D points P (Nx3)
95
96     ***
97     Hints:
98     (1) Loop through the 'M2s' and use triangulate to calculate the 3D points and projection error. Keep track
99         of the projection error through best_error and retain the best one.
100    (2) Remember to take a look at camera2 to see how to correctly retrieve the M2 matrix from 'M2s'.
101
102    ***
103    # ----- TODO -----
104    # YOUR CODE HERE
105    # Calculate E from F and K1/K2
106    K1, K2 = intrinsics["K1"], intrinsics["K2"]
107    E = essentialMatrix(F, K1, K2)
108
109    # Calculate M1 & M2
110    M1 = np.concatenate([np.eye(3), np.zeros((3, 1))], axis=1) #[I|0]
111    M2s = camera2(E)
112    C1 = K1 @ M1
113
114    # Select the M2 that has most positive depth points
115    max_posdepth_num = -1
116    min_error = np.inf
117    M2_ret = M2s[:, :, 0]
118    P_ret = np.zeros((pts1.shape[0], 3))
119    for index in range(M2s.shape[-1]):
120        M2 = M2s[:, :, index]
121        C2 = K2 @ M2
122        P, error = triangulate(C1, pts1, C2, pts2)
123
124        ...
125
126        if error < min_error:
127            min_error = error
128            M2_ret = M2.copy()
129            P_ret = P.copy()
129
130        posdepth_num = np.sum((P[:, 2] > 0))
131        if posdepth_num > max_posdepth_num:
132            max_posdepth_num = posdepth_num
133            M2_ret = M2.copy()
134            P_ret = P.copy()
135
136    C2_ret = K2 @ M2_ret
137
138    # Write file to q3_3.npz
139    np.savez(filename, M2_ret, C2_ret, P_ret)
140    return M2_ret, C2_ret, P_ret

```

Figure 15: Code Snippet

4 3D Visualization

You will now create a 3D visualization of the temple images. By treating our two images as a stereo-pair, we can triangulate corresponding points in each image, and render their 3D locations.

Q4.1 [15 points] In `q4_1_epipolar_correspondence.py` finish the function `epipolarCorrespondence` with the signature:

```
[x2, y2] = epipolarCorrespondence(im1, im2, F, x1, y1)
```

This function takes in the x and y coordinates of a pixel on `im1` and your fundamental matrix `F`, and returns the coordinates of the pixel on `im2` which correspond to the input point. The match is obtained by computing the similarity of a small window around the (x_1, y_1) coordinates in `im1` to various windows around possible matches in the `im2` and returning the closest.

Instead of searching for the matching point at every possible location in `im2`, we can use `F` and simply search over the set of pixels that lie along the epipolar line (recall that the epipolar line passes through a single point in `im2` which corresponds to the point (x_1, y_1) in `im1`).

There are various possible ways to compute the window similarity. For this assignment, simple methods such as the Euclidean or Manhattan distances between the intensity of the pixels should suffice. See [2] chapter 11, on stereo matching, for a brief overview of these and other methods.

Implementation hints:

- Experiment with various window sizes.
- It may help to use a Gaussian weighting of the window, so that the center has greater influence than the periphery.
- Since the two images only differ by a small amount, it might be beneficial to consider matches for which the distance from (x_1, y_1) to (x_2, y_2) is small.

To help you test your `epipolarCorrespondence`, we have included a helper function `epipolarMatchGUI` in `q4_1_epipolar_correspondence.py`, which takes in two images and the fundamental matrix. This GUI allows you to click on a point in `im1`, and will use your function to display the corresponding point in `im2`. See [Figure 16](#).

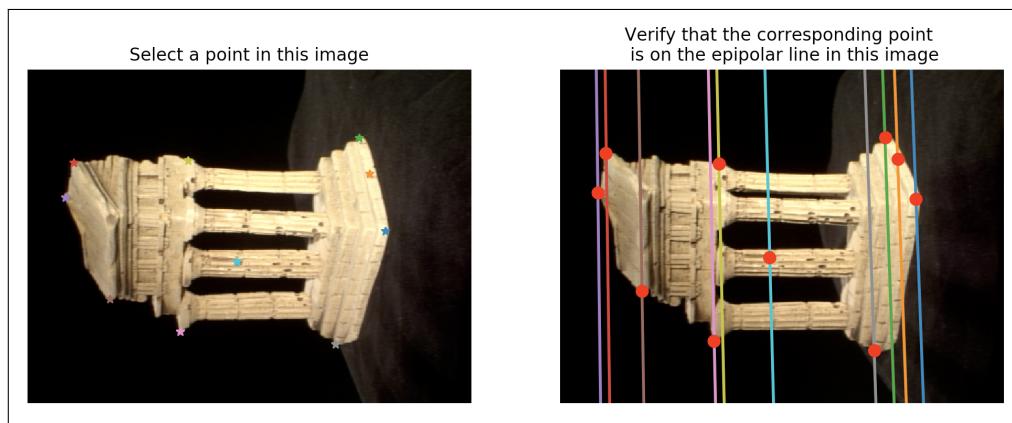


Figure 16: `epipolarMatchGUI` shows the corresponding point found by calling `epipolarCorrespondence`

It's not necessary for your matcher to get *every* possible point right, but it should get easy points (such as those with distinctive, corner-like windows). It should also be good enough to render an intelligible representation in the next question.

Output: Save the matrix **F**, points **pts1** and **pts2** which you used to generate the screenshot to the file **q4_1.npz**.

In your write-up:

- Include a screenshot of `epipolarMatchGUI` with some detected correspondences.
- Include the code snippet of `epipolarCorrespondence` function.

Q4.1

The following **Figure 17** shows the screenshot of `epipolarMatchGUI()` with some detected correspondences using the `epipolarCorrespondence()` function:

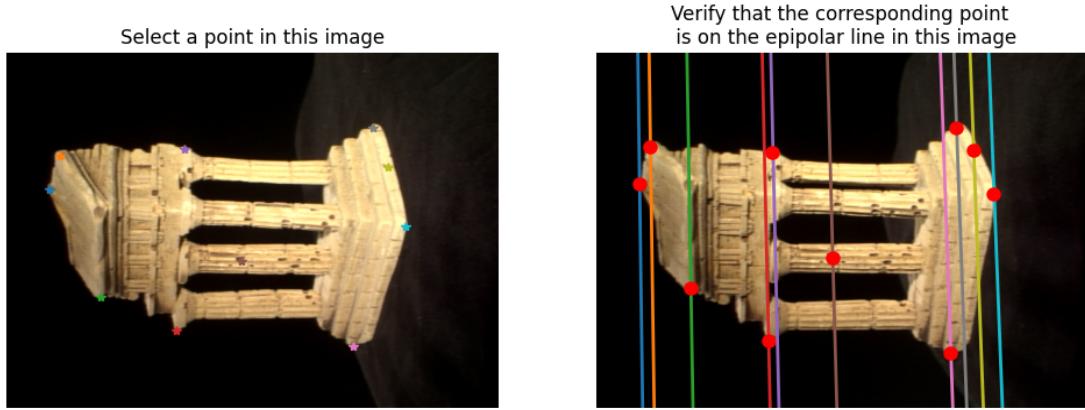


Figure 17: Result of the `epipolarCorrespondence()` Function

Q4.1 continued

The following [Figure 18](#) shows the code snippet of the epipolarCorrespondence() function in q4_1_epipolar_correspondence.py:

```

182 def epipolarCorrespondence(im1, im2, F, x1, y1):
183     # Replace pass by your implementation
184     # ----- TODO -----
185     # YOUR CODE HERE
186     # Construct l2 line equation
187     x1_pt = np.array([x1, y1, 1]).reshape(-1, 1)
188     l2_line = F @ x1_pt
189
190     # Extract points in im2 along the l2 line
191     x2_pts = extractPointsAlongLine(l2_line, im_width=im2.shape[1], im_height=im2.shape[0])
192
193     # Extract x1_roi
194     ksize = 71
195     #print(f"-----({x1, y1}) = {(x1, y1)}")
196     x1_roi = extractROI(image=im1, x=x1, y=y1, ksize=ksize)
197
198     # Get the best pt in im2
199     best_err = np.inf
200     best_cand = (None, None)
201     for x2_pt_candi in x2_pts:
202         #print(f"====")
203         # Extract x2_roi
204         x2_roi = extractROI(image=im2, x=x2_pt_candi[0], y=x2_pt_candi[1], ksize=ksize)
205
206         # Apply Gaussian weighting
207         gaussien2DKernel = getGaussianWeight(ksize=ksize, sigma=5.0, print_verbose=False)
208         #x1_gauss = applyGaussianKernel(x1_roi, gaussian2DKernel=gaussien2DKernel)
209         #x2_gauss = applyGaussianKernel(x2_roi, gaussian2DKernel=gaussien2DKernel)
210
211         # Calculate the similarity
212         err = calculateError(x1_roi=x1_roi, x2_roi=x2_roi, kernel=gaussien2DKernel)
213         #print(f"x2_pt_candi, err = {x2_pt_candi}, {err}")
214
215         # Get the best
216         if err < best_err:
217             best_err = err
218             best_cand = x2_pt_candi
219
220     #x2, y2 = best_cand
221     #error = np.linalg.norm(np.array([x2, y2]) - np.array([x1, y1]))
222     #print(f"error = {error}")
223     #print(f"best_cand = {best_cand}")
224     return best_cand

```

Figure 18: Code Snippet of the epipolarCorrespondence() Function

Q4.2 [10 points] Included in this homework is a file `data/templeCoords.npz` which contains 288

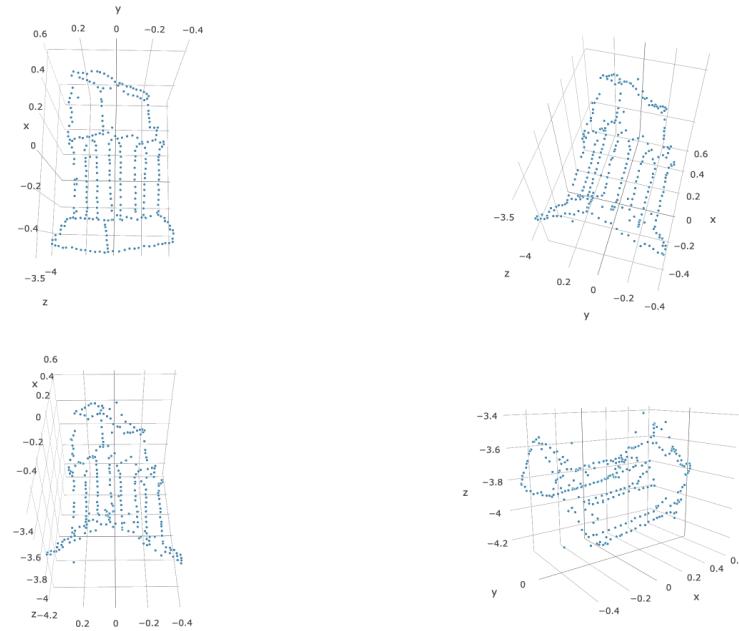


Figure 19: An example point cloud

hand-selected points from `im1` saved in the variables `x1` and `y1`.

Now, we can determine the 3D location of these point correspondences using the `triangulate` function. These 3D point locations can then be plotted using the Matplotlib or plotly package. Complete the `compute3D_pts` function in `q4_2_visualize.py`, which loads the necessary files from `..../data/` to generate the 3D reconstruction using `scatter` function matplotlib. An example is shown in [Figure 19](#).

Output: Again, save the matrix **F**, matrices **M1**, **M2**, **C1**, **C2** which you used to generate the screenshots to the file `q4_2.npz`.

In your write-up:

- Take a few screenshots of the 3D visualization so that the outline of the temple is clearly visible, and include them in your writeup.
- Include the code snippet of `compute3D_pts` function in your write-up.

Q4.2

The following [Figure 20](#), [Figure 21](#), [Figure 22](#) and [Figure 23](#) show the results of the 3D visualization results of the temple:

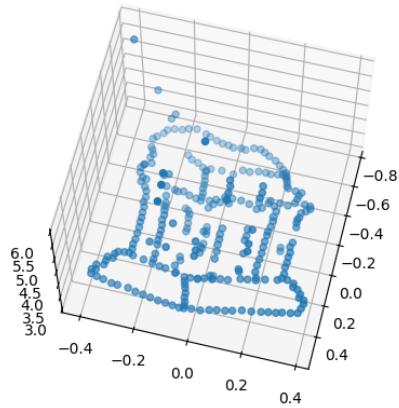


Figure 20: First Screenshot

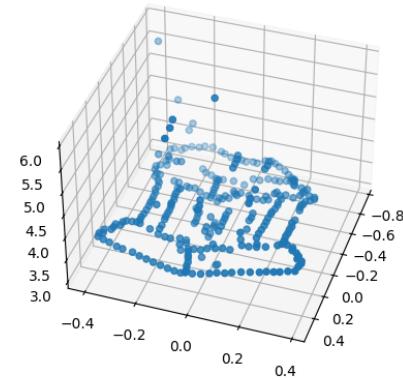


Figure 21: Second Screenshot

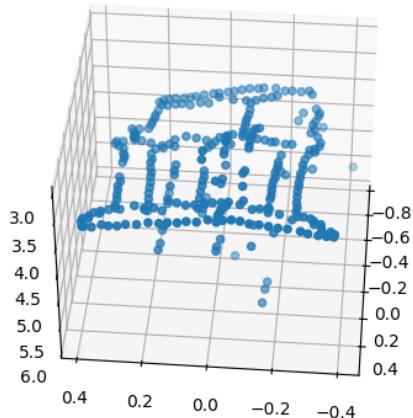


Figure 22: Third Screenshot

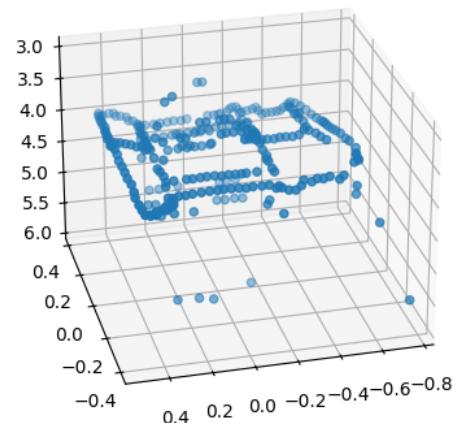


Figure 23: Fourth Screenshot

Q4.2 continued

The following [Figure 24](#) shows the code snippet of the `compute3D_pts()` function in `q4_2_visualize.py`:

```

17 """
18 Q4.2: Finding the 3D position of given points based on epipolar correspondence and triangulation
19     Input: temple_pts1, chosen points from im1
20             intrinsics, the intrinsics dictionary for calling epipolarCorrespondence
21             F, the fundamental matrix
22             im1, the first image
23             im2, the second image
24     Output: P (Nx3) the recovered 3D points
25
26 Hints:
27 (1) Use epipolarCorrespondence to find the corresponding point for [x1 y1] (find [x2, y2])
28 (2) Now you have a set of corresponding points [x1, y1] and [x2, y2], you can compute the M2
29     matrix and use triangulate to find the 3D points.
30 (3) Use the function findM2 to find the 3D points P (do not recalculate fundamental matrices)
31 (4) As a reference, our solution's best error is around ~2200 on the 3D points.
32
33 Modified by Vineet Tambe, 2023.
34 """
35
36
37 def compute3D_pts(temple_pts1, intrinsics, F, im1, im2, filename='q4_2.npz'):
38     # ----- TODO -----
39     # YOUR CODE HERE
40     # Calculate M1 & C1
41     K1, K2 = intrinsics["K1"], intrinsics["K2"]
42     M1 = np.concatenate([np.eye(3), np.zeros((3, 1))], axis=1) #[I|0]
43     C1 = K1 @ M1
44     # Get corresponding point [x2, y2]
45     temple_pts2 = np.zeros(temple_pts1.shape, dtype=int)
46     for index, (x1, y1) in enumerate(temple_pts1):
47         temple_pts2[index] = epipolarCorrespondence(im1, im2, F, x1, y1)
48
49     # Compute M2 matrix
50     M2, C2, P = findM2(F, temple_pts1, temple_pts2, intrinsics, filename="q4_2_3.npz")
51
52     # Write file to filename
53     np.savez(filename, F, M1, M2, C1, C2)
54     return P

```

Figure 24: Code Snippet of the `compute3D_pts()` Function

5 Bundle Adjustment

Bundle Adjustment is commonly used as the last step of every feature-based 3D reconstruction algorithm. Given a set of images depicting a number of 3D points from different viewpoints, bundle adjustment is the process of simultaneously refining the 3D coordinates along with the camera parameters. It minimizes reprojection error, which is the squared sum of distances between image points and predicted points. In this section, you will implement bundle adjustment algorithm by yourself (make use of `q5_bundle_adjustment.py` file). Specifically,

- In Q5.1, you need to implement a RANSAC algorithm to estimate the fundamental matrix \mathbf{F} and all the inliers.
- In Q5.2, you will need to write code to parameterize Rotation matrix \mathbf{R} using [Rodrigues formula](#) (please check [this pdf](#) for a detailed explanation), which will enable the joint optimization process for Bundle Adjustment.
- In Q5.3, you will need to first write down the objective function in `rodriguesResidual`, and do the `bundleAdjustment`.

Q5.1 RANSAC for Fundamental Matrix Recovery [15 points] In some real world applications, manually determining correspondences is infeasible and often there will be noisy correspondences. Fortunately, the RANSAC method seen in class can be applied to the problem of fundamental matrix estimation.

Implement the above algorithm with the signature:

```
[F, inliers] = ransacF(pts1, pts2, M, nIters, tol)
```

where M is defined in the same way as in [section 2](#) and `inliers` is a boolean vector of size equivalent to the number of points. Here `inliers` is set to true only for the points that satisfy the threshold defined for the given fundamental matrix \mathbf{F} .

We have provided some noisy correspondences in `some_corresp_noisy.npz` in which around 75% of the points are inliers.

In your write-up: Compare the result of RANSAC with the result of the eightpoint when ran on the noisy correspondences. Briefly explain the error metrics you used, how you decided which points were inliers, and any other optimizations you may have made. `nIters` is the maximum number of iterations of RANSAC and `tol` is the tolerance of the error to be considered as inliers. Discuss the effect on the Fundamental matrix by varying these values. **Please include the code snippet of the `ransacF` function in your write-up.**

- *Hints:* Use the Eight or Seven point algorithm to compute the fundamental matrix from the minimal set of points. Then compute the inliers, and refine your estimate using all the inliers.

Q5.1

The following **Figure 25** shows the screenshot of RANSAC Seven Point algorithm result:

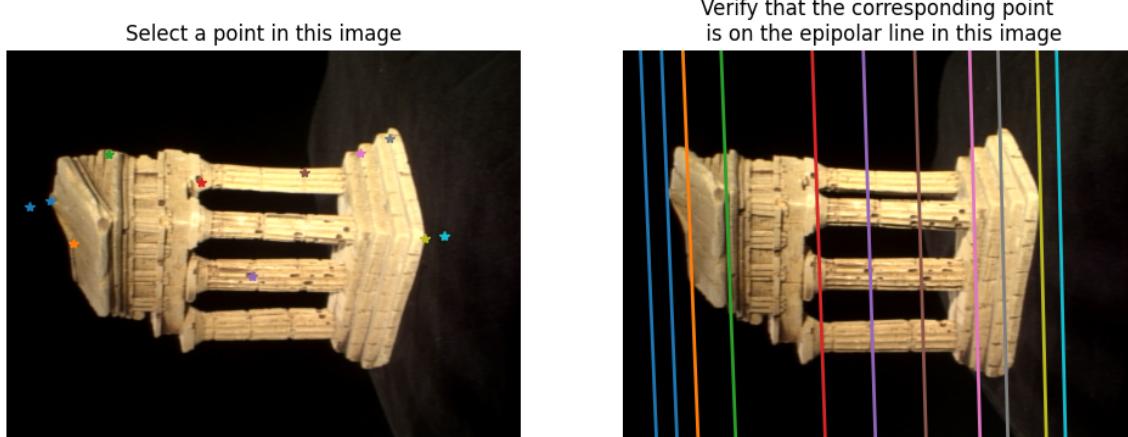


Figure 25: Result of the RANSAC Seven Point Algorithm

The following **Figure 26** shows the screenshot of Eight Point algorithm result:

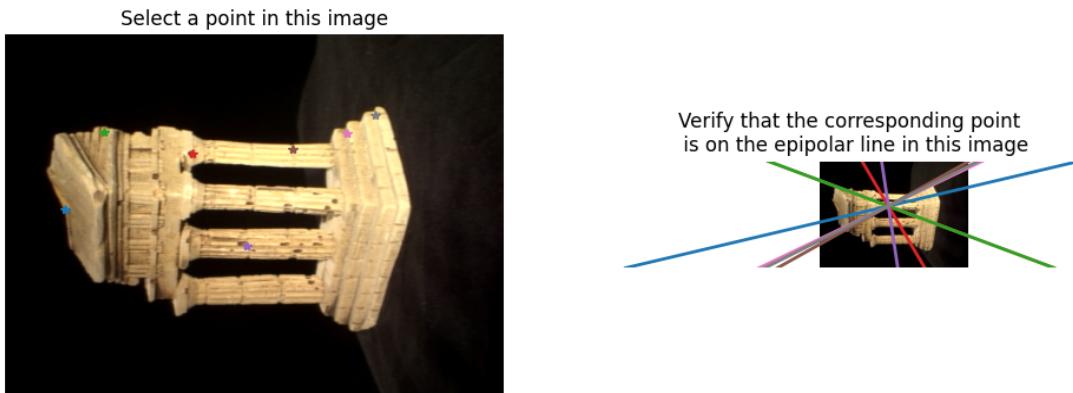


Figure 26: Result of the Eight Point Algorithm

In RANSAC Seven Point algorithm, I use `nIters=1000, tol=2` to get above result, and the inlier number is 106, which is around 75% with total number of points is 140. The reason that the RANSAC Seven Point algorithm gets better result is that the points are noisy, and it is stated in the hw3.pdf that around 75% of these noisy points are inliers. While RANSAC Seven Point algorithm has the chance to reject outliers, the Eight Point algorithm can only use all the points, including outliers, to calculate the Fundamental Matrix, F. Accordingly, RANSAC Seven Point algorithm performs better than Eight Point algorithm under noisy data.

Q5.1 continued

The following shows the results of changing the value of the 'tol' parameter.

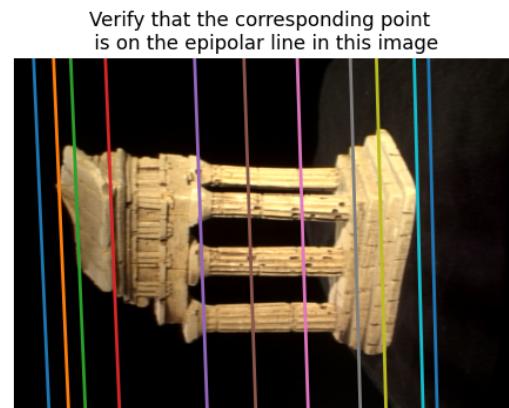
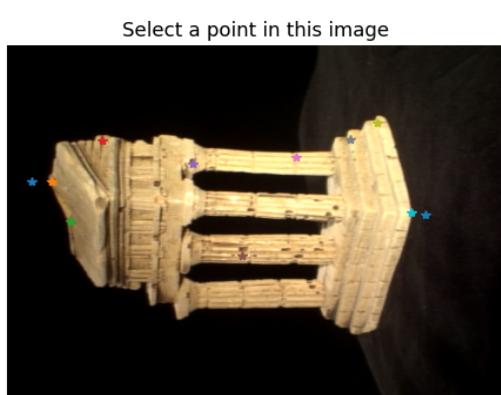


Figure 27: Result of the RANSAC Seven Point Algorithm w/ tol=0.5, inlier rate=52.86%

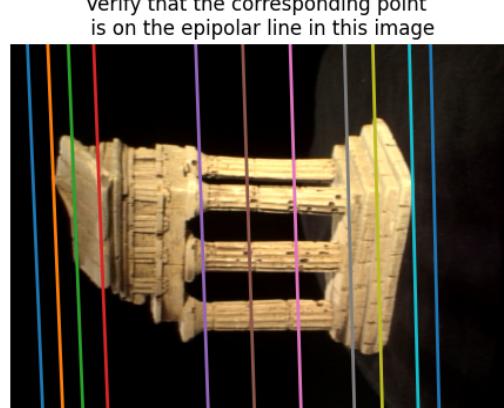
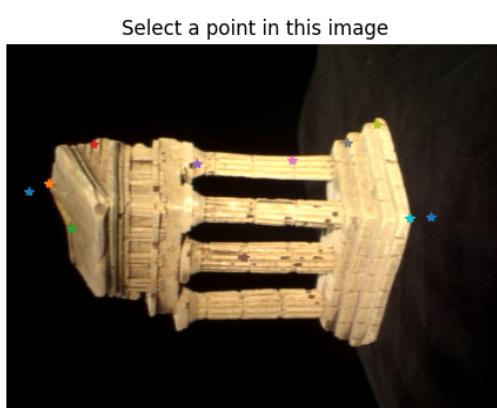


Figure 28: Result of the RANSAC Seven Point Algorithm w/ tol=1, inlier rate=65.71%

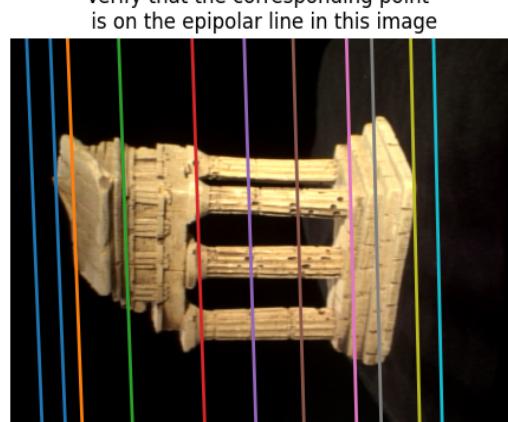
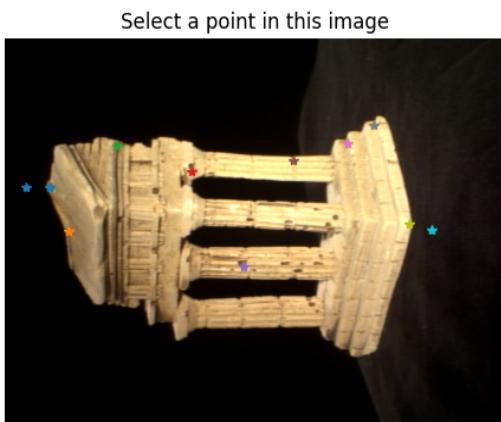


Figure 29: Result of the RANSAC Seven Point Algorithm w/ tol=4, inlier rate=76.43%

Q5.1 continued

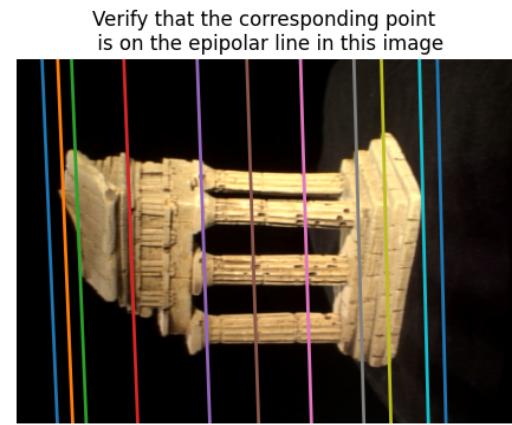
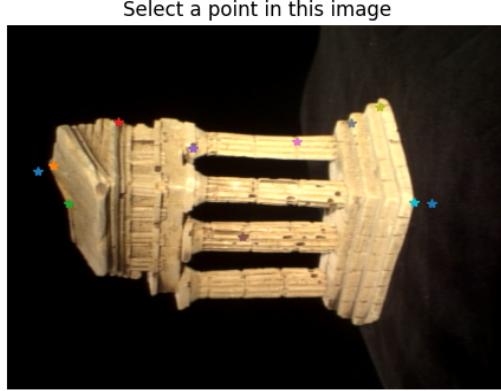


Figure 30: Result of the RANSAC Seven Point Algorithm w/ tol=8, inlier rate=78.57%

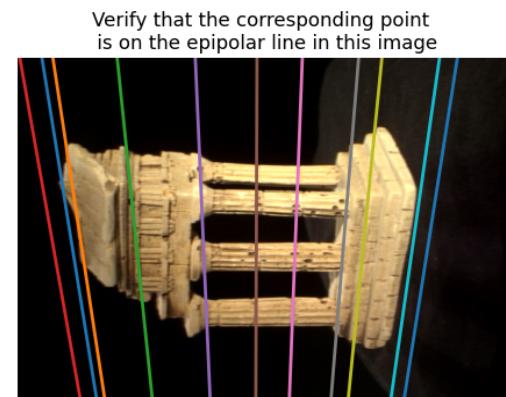
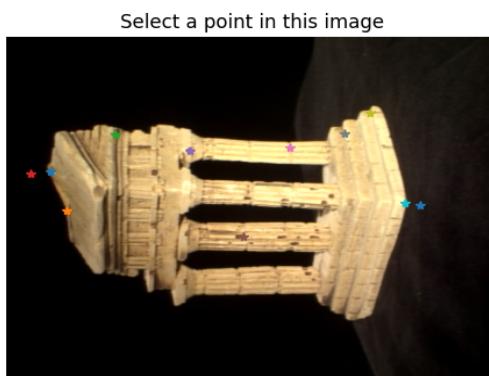


Figure 31: Result of the RANSAC Seven Point Algorithm w/ tol=16, inlier rate=79.29%

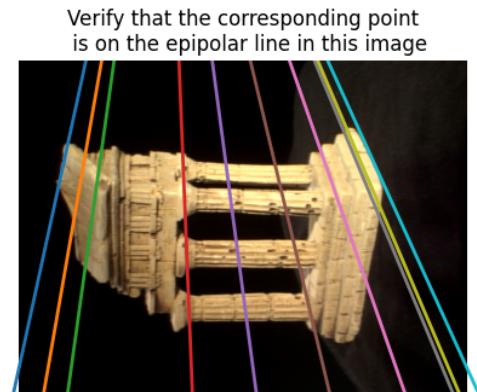
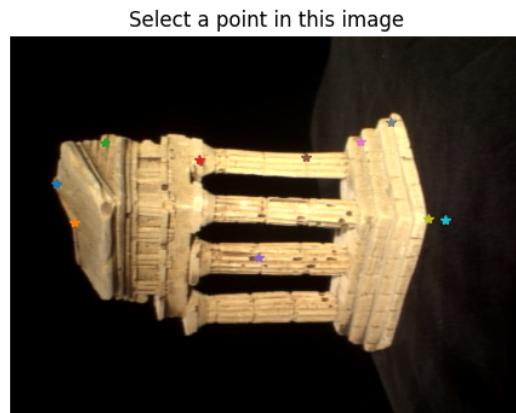


Figure 32: Result of the RANSAC Seven Point Algorithm w/ tol=128, inlier rate=80%

Q5.1 continued

As we can see from the results above that when tol is getting larger, specifically larger than 8, the result is getting worse. It is because the criteria to count a point as an inlier is getting laxer, and some of the outlier points are wrongly categorized as inliers. The resulted F is generated by some outliers and thus becomes worse. As for the necessary iteration number of RANSAC to return a good model, the value is calculated by $k = \frac{\log(1-p)}{\log(1-w^n)}$, where $p = 0.999$, w is about 0.75 (from hw3.pdf), and $n=7$, then $k = 48.21$, so the nIters=1000, as I set in the program for above results, should be sufficient to have RANSAC Seven Point algorithm return a good model.

The following [Figure 33](#) shows the code snippet of the ransacF() function in q5_bundle_adjustment.py:

```

34 """
35 Q5.1: RANSAC method.
36   Input: pts1, Nx2 Matrix
37   pts2, Nx2 Matrix
38   M, a scalar parameter
39   nIters, Number of iterations of the Ransac
40   tol, tolerance for inliers
41   Output: F, the fundamental matrix
42           inliers, Nx1 bool vector set to true for inliers
43
44 Hints:
45 (1) You can use the calc_epi_error from q1 with threshold to calculate inliers. Tune the threshold based on
46     the results/expected number of inliners. You can also define your own metric.
47 (2) Use the seven point algorithm to estimate the fundamental matrix as done in q1
48 (3) Choose the resulting F that has the most number of inliers
49 (4) You can increase the nIters to bigger/smaller values
50
51 """
52 def getFSevenPoint(pts1_choice, pts2_choice, pts1_homo, pts2_homo, M, tol, file="debug.log"):
53     best_inlier = np.zeros((pts1_homo.shape[0], 1), dtype=bool)
54     best_F = None
55     best_inlier_cnt = -1
56
57     Fs = sevenpoint(pts1_choice, pts2_choice, M)
58     for F in Fs:
59         inlier = np.zeros((pts1_homo.shape[0], 1), dtype=bool)
60         error_all = np.abs(calc_epi_error(pts1_homo, pts2_homo, F))
61         inlier[error_all < tol] = True
62         inlier_cnt = np.sum(inlier)
63
64         if inlier_cnt > best_inlier_cnt:
65             best_inlier_cnt = inlier_cnt
66             best_F = F.copy()
67             best_inlier = inlier.copy()
68
69     return best_F, best_inlier, best_inlier_cnt
70
71 def ransacF(pts1, pts2, M, nIters=1000, tol=10):
72     # TODO: Replace pass by your implementation
73     pts1_homo = np.hstack((pts1, np.ones((pts1.shape[0], 1))))
74     pts2_homo = np.hstack((pts2, np.ones((pts2.shape[0], 1))))
75     best_inlier = np.zeros((pts1_homo.shape[0], 1), dtype=bool)
76     best_F = None
77     best_inlier_cnt = -1
78
79     file = open("debug.log", "w")
80     for i in range(nIters):
81         inlier = np.zeros((pts1.shape[0], 1), dtype=bool)
82         choice = np.random.choice(range(pts1.shape[0]), 7)
83         pts1_choice = pts1[choice, :]
84         pts2_choice = pts2[choice, :]
85         cur_F, cur_inlier, cur_inlier_cnt = getFSevenPoint(pts1_choice, pts2_choice, pts1_homo, pts2_homo, M, tol, file=file)
86
87         if cur_inlier_cnt > best_inlier_cnt:
88             best_inlier_cnt = cur_inlier_cnt
89             best_F = cur_F.copy()
90             best_inlier = cur_inlier.copy()
91
92     return best_F, best_inlier

```

[Figure 33](#): Code Snippet of the ransacF() Function

Q5.2 Rodrigues and Invsere Rodrigues [15 points] So far we have independently solved for camera matrix, \mathbf{M}_j and 3D points \mathbf{w}_i . In bundle adjustment, we will jointly optimize the reprojection error with respect to the points \mathbf{w}_i and the camera matrix \mathbf{C}_j .

$$err = \sum_{ij} \|\mathbf{x}_{ij} - Proj(\mathbf{C}_j, \mathbf{w}_i)\|^2,$$

where $\mathbf{C}_j = \mathbf{K}_j \mathbf{M}_j$, same as in Q3.2.

For this homework we are going to only look at optimizing the extrinsic matrix. To do this we will be parameterizing the rotation matrix \mathbf{R} using Rodrigues formula to produce vector $\mathbf{r} \in \mathbb{R}^3$. Write a function that converts a Rodrigues vector \mathbf{r} to a rotation matrix \mathbf{R}

$$\mathbf{R} = \text{rodrigues}(\mathbf{r})$$

as well as the inverse function that converts a rotation matrix \mathbf{R} to a Rodrigues vector \mathbf{r}

$$\mathbf{r} = \text{invRodrigues}(\mathbf{R})$$

Reference: [Rodrigues formula](#) and [this pdf](#).

In your write-up: Include the code snippet of `rodrigues` and `invRodrigues` functions.

Q5.2

The following [Figure 34](#), [Figure 35](#) show the code snippet of `rodrigues()` and `invRodrigues()` functions respectively:

```

94 """
95 Q5.2: Rodrigues formula.
96   Input: r, a 3x1 vector
97   Output: R, a rotation matrix
98 """
99
100
101 def rodrigues(r, epsilon=1e-20):
102     # TODO: Replace pass by your implementation
103     r_reshape = r.reshape((-1, 1))
104     theta = np.linalg.norm(r_reshape) # floating point value
105     if theta < epsilon: # == 0
106         return np.eye(3)
107
108     u = r_reshape/theta
109     u_cross = np.array([[ 0, -u[2, 0], u[1, 0]],
110                         [ u[2, 0], 0, -u[0, 0]],
111                         [-u[1, 0], u[0, 0], 0]])
112     R = np.eye(3)*np.cos(theta) + (1-np.cos(theta))*(u@u.T) + u_cross*np.sin(theta)
113     return R
114

```

Figure 34: Code Snippet of `rodrigues()`.

```

111 """
112   Q5.2: Inverse Rodrigues formula.
113   Input: R, a rotation matrix
114   Output: r, a 3x1 vector
115 """
116
117 def invRodrigues(r, epsilon=1e-10):
118     # TODO: Replace pass by your implementation
119     a = input()
120     if a == None:
121         raise ValueError("Input must be np.pi and (r[0, 0] < epsilon and r[1, 0] < epsilon and r[2, 0] < 0) or (r[0, 0] < epsilon and r[1, 0] > 0) or (r[0, 0] > 0);")
122     else:
123         return a
124
125     return r
126
127 def invRodrigues(r, epsilon=1e-10):
128     # TODO: Replace pass by your implementation
129     a = (0, 0, 0)
130     A = np.array([A[0, 1], A[0, 2], A[1, 1]]).T
131     A = A - np.linalg.norm(A)
132     c = (A[0, 0] + A[1, 1] + A[2, 2]) / 3
133
134     if c < epsilon and (c+1) < epsilon: # if c == 0 and c == 1
135         a = np.array([0, 0, 0]).T
136     elif c < epsilon and (c+1) > epsilon: # if c == 0 and c == -1
137         a = np.array([0, 0, 0]).T
138     else:
139         matrix = r + r * np.eye(3)
140         for i in range(matrix.shape[1]):
141             if np.any(matrix[:, i]): # Check if any element in the column is non-zero
142                 a[0, i] = matrix[0, i]
143                 break
144
145         a[0, 0] = np.linalg.norm(r)
146
147         return a
148     else:
149         w = np.linalg.norm(r)
150         if w != 0:
151             theta = np.arctan2(w, c)
152             matrix = np.eye(3) - r * np.reshape(theta, (3, 1))
153             return np.linalg.inv(matrix) @ r
154
155     raise ValueError("Unpredicted combination of values of a and c.")

```

Figure 35: Code Snippet of `invRodrigues()`.

Q5.3 Bundle Adjustment [10 points]

Using this parameterization, write an optimization function

$$\text{residuals} = \text{rodriguesResidual}(K_1, M_1, p_1, K_2, p_2, x)$$

where x is the flattened concatenation of \mathbf{x} , \mathbf{r}_2 , and \mathbf{t}_2 . \mathbf{w} are the 3D points; \mathbf{r}_2 and \mathbf{t}_2 are the rotation (in the Rodrigues vector form) and translation vectors associated with the projection matrix \mathbf{M}_2 . The `residuals` are the difference between original image projections and estimated projections (the square of $L2$ -norm of this vector corresponds to the error we computed in Q3.2):

$$\text{residuals} = \text{numpy.concatenate}([(p_1 - p_1_\hat{ }) . \text{reshape}([-1]), (p_2 - p_2_\hat{ }) . \text{reshape}([-1])])$$

Use this error function and Scipy's optimizer `minimize` to write a function that optimizes for the best extrinsic matrix and 3D points using the inlier correspondences from `some_corresp_noisy.npz` and the RANSAC estimate of the extrinsics and 3D points as an initialization.

```
[M2, w, o1, o2] = bundleAdjustment(K1, M1, p1, K2, M2_init, p2, w_init)
```

Try to extract the rotation and translation from `M2_init`, then use `invRodrigues` you implemented previously to transform the rotation, concatenate it with translation and the 3D points, then the concatenated vector are variables to be optimized. After obtaining optimized vector, decompose it back to rotation using `rodrigues` you implemented previously, translation and 3D points coordinates.

In your write-up:

- Include an image of the original 3D points and the optimized points (use the provided `plot_3D_dual` function).
- Report the reprojection error with your initial M_2 and w , as well as with the optimized matrices.
- Include the code snippets for `rodriguesResidual` and `bundleAdjustment` in your write-up.

Hint: For reference, our solution achieves a reprojection error around 10 after optimization. Your exact error may differ slightly.

Q5.3

The following [Figure 36](#), [Figure 37](#), [Figure 38](#), and [Figure 39](#) show the results of the optimization.

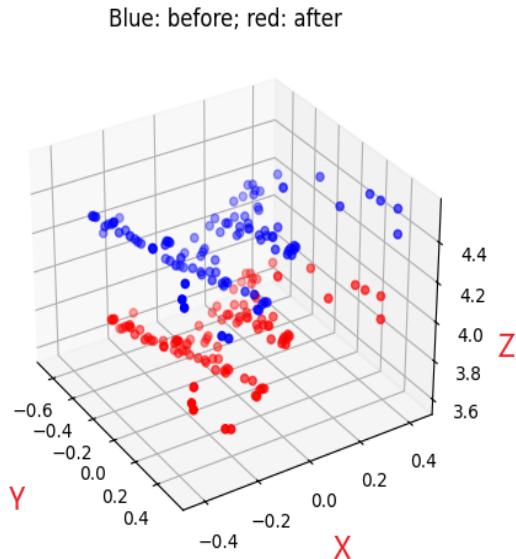


Figure 36: Perspective 1 of the Result of Optimization

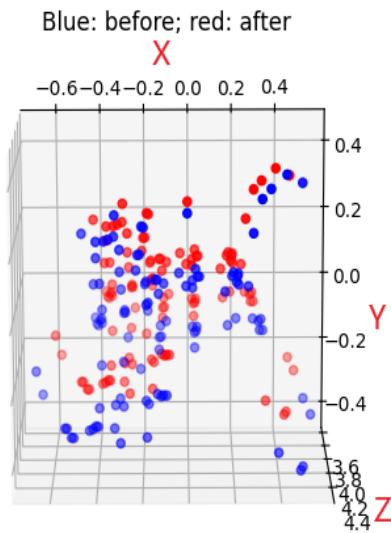


Figure 37: Perspective 2 of the Result of Optimization

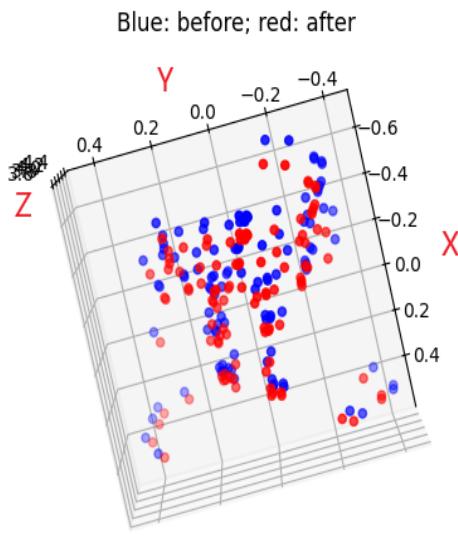


Figure 38: Perspective 3 of the Result of Optimization

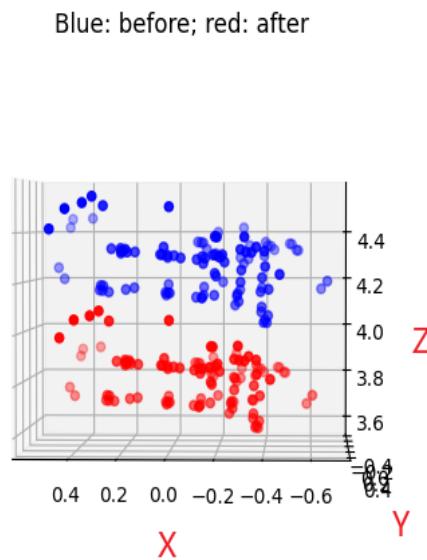


Figure 39: Perspective 4 of the Result of Optimization

Q5.3 continued

The reprojection error with initial M2 and w is 874.0899779851011.
The reprojection error with optimized M2 and w is 8.08757938434745.
The following [Figure 40](#), [Figure 41](#), and [Figure 42](#) show the code snippets.

```

311 # Visualization:
312 np.random.seed(1)
313 correspondence = np.load(
314     "data/some_corresp_noisy.npz"
315 ) # Loading noisy correspondences
316 intrinsics = np.load("data/intrinsics.npz") # Loading the intrinsics of the camera
317 K1, K2 = intrinsics["K1"], intrinsics["K2"]
318 pts1, pts2 = correspondence["pts1"], correspondence["pts2"]
319 im1 = plt.imread("data/im1.png")
320 im2 = plt.imread("data/im2.png")
321 M = np.max([*im1.shape, *im2.shape])
322
323 # TODO: YOUR CODE HERE
324 """
325 Call the ransacF function to find the fundamental matrix
326 Call the findM2 function to find the extrinsics of the second camera
327 Call the bundleAdjustment function to optimize the extrinsics and 3D points
328 Plot the 3D points before and after bundle adjustment using the plot_3D_dual function
329 """
330
331 #Step1: Call the ransacF function
332 F, inliers = ransacF(pts1, pts2, M=np.max([*im1.shape, *im2.shape]), nIters=1000, tol=2)
333 print(f"number of inliers = {np.sum(inliers)}\ntotal number of points = {pts1.shape[0]}\ninlier rate = {np.sum(inliers)/pts1.shape[0]*100}%")
334
335 #displayEpipolarF(im1, im2, F)
336 np.savez('q5_3.npz', F, inliers)
337 #F = np.load('q5_3.npz')['arr_0']
338 #inliers = np.load('q5_3.npz')['arr_1']
339 #displayEpipolarF(im1, im2, F)
340 #print(f"F = {F}")
341
342 #Step2: Call the findM2 function to find the M (extrinsics) of the second camera
343 p1, p2 = pts1[inliers.flatten()], pts2[inliers.flatten()]
344 M2_init, C2_init, P_init = findM2(F, p1, p2, intrinsics, filename="q5_3_findM2.npz")
345 M1 = np.concatenate([np.eye(3), np.zeros((3, 1))], axis=1) #[1|0]
346
347 #Step3: Call the bundleAdjustment function to optimize the extrinsics and 3D points
348 M2_opt, P_opt, object_start, object_end = bundleAdjustment(K1, M1, p1, K2, M2_init, p2, P_init)
349 print(f"Reprojection error with initial M2 and w = {object_start}")
350 print(f"Optimized Reprojection error with optimized M2 and w = {object_end}")
351
352 #Step4: Plot the 3D points before and after bundle adjustment using the plot_3D_dual function
353 fig = plt.figure()
354 ax = Axes3D(fig)
355 plot_3D_dual(P_init, P_opt)

```

Figure 40: Code Snippet of Main (Test) Function.

Q5.3 continued

```

217 """
218 Q5.3 Bundle adjustment.
219 Input: K1, the intrinsics of camera 1
220 M1, the extrinsics of camera 1
221 p1, the 2D coordinates of points in image 1
222 K2, the intrinsics of camera 2
223 M2_init, the initial extrinsics of camera 1
224 p2, the 2D coordinates of points in image 2
225 P_init, the initial 3D coordinates of points
226 Output: M2, the optimized extrinsics of camera 1
227 P, the optimized 3D coordinates of points
228 o1, the starting objective function value with the initial input
229 o2, the ending objective function value after bundle adjustment
230

231 Hints:
232 (1) Use the scipy.optimize.minimize function to minimize the objective function, rodriguesResidual.
233 You can try different (method='...') in scipy.optimize.minimize for best results.
234 """
235
236 def bundleAdjustment(K1, M1, p1, K2, M2_init, p2, P_init):
237     """
238         YOUR CODE HERE
239     # Extract r2, t2 from M2_init
240     R2 = M2_init[0:3, 0:3]
241     r2 = invRodrigues(R2, 1e-32).reshape(-1, 1) #3x1
242     t2 = M2_init[3:, 0].reshape(-1, 1) #3x1
243     init_params = np.concatenate([P_init.flatten(), r2.flatten(), t2.flatten()])
244
245     # Calculate the starting object loss
246     obj_start = np.sum(rodriguesResidual(K1, M1, p1, K2, p2, init_params)**2)
247
248     # Optimization
249     object_func = lambda params: np.sum(rodriguesResidual(K1, M1, p1, K2, p2, params)**2)
250     result = scipy.optimize.minimize(
251         object_func,          # Objective function (residual function)
252         fun=object_func,      # Initial guess for parameters
253         x0=init_params,       # Initial guess for parameters
254         method='L-BFGS-B',    # Optimization method (you can try different methods)
255     )
256     optimized_params = result.x
257
258     # Extract optimized r2, t2, P, and M2
259     t2_optimized = optimized_params[-3:].reshape(-1, 1) #3x1
260     r2_optimized = optimized_params[6:-3].reshape(-1, 1) #3x1
261     P_optimized = optimized_params[0:6].reshape(-1, 1) #3x3
262     R2_optimized = invRodrigues(r2_optimized, 1e-32).reshape(-1, 1) #3x3
263     M2 = np.hstack((R2_optimized, t2_optimized))
264
265     # Calculate the resulting object loss
266     obj_end = np.sum(rodriguesResidual(K1, M1, p1, K2, p2, optimized_params)**2)
267
268     return M2, P_optimized, obj_start, obj_end
269

```

Figure 41: Code Snippet of bundleAdjustment()

```

173 """
174 Q5.3: Rodrigues residual.
175 Input: K1, the intrinsics of camera 1
176 M1, the extrinsics of camera 1
177 p1, the 2D coordinates of points in image 1
178 K2, the intrinsics of camera 2
179 p2, the 2D coordinates of points in image 2
180 x, the flattened concatenationg of P, r2, and t2.
181 Output: residuals, 4N x 1 vector, the difference between original and estimated projections
182 """
183
184 def rodriguesResidual(K1, M1, p1, K2, p2, x):
185     """
186         # TODO: Replace pass by your implementation
187         P = x[:6].reshape(-1, 3) #Nx3
188         r2 = x[6:9].reshape(-1, 1) #Nx3
189         t2 = x[9:12].reshape(-1, 1) #Nx3
190
191     # Reconstruct C1
192     C1 = K1@M1
193
194     # Reconstruct C2
195     R2 = rodrigues(r2)
196     M2 = np.hstack((R2, t2))
197     C2 = K2@M2
198
199     # Calculate projected points on C1 and C2
200     P_homo = np.hstack((P, np.ones((P.shape[0], 1)))) #Nx4
201     p1_proj_homo = C1@P_homo[:, 1] #Nx3
202     p2_proj_homo = C2@P_homo[:, 1] #Nx3
203
204     # Convert Homogeneous back to normal coordinates
205     p1_hat_x = p1_proj_homo[:, 0]/p1_proj_homo[:, 2]
206     p1_hat_y = p1_proj_homo[:, 1]/p1_proj_homo[:, 2]
207     p1_hat_z = np.column_stack((p1_hat_x, p1_hat_y)) #Nx2
208     p2_hat_x = p2_proj_homo[:, 0]/p2_proj_homo[:, 2]
209     p2_hat_y = p2_proj_homo[:, 1]/p2_proj_homo[:, 2]
210     p2_hat_z = np.column_stack((p2_hat_x, p2_hat_y)) #Nx2
211
212     # Calculate the residual 4Nx1 vector
213     residual = np.concatenate([(p1_hat_x - p1).reshape(-1), (p2_hat_y - p2).reshape(-1)])
214
215     return residual
216

```

Figure 42: Code Snippet of rodriguesResidual()

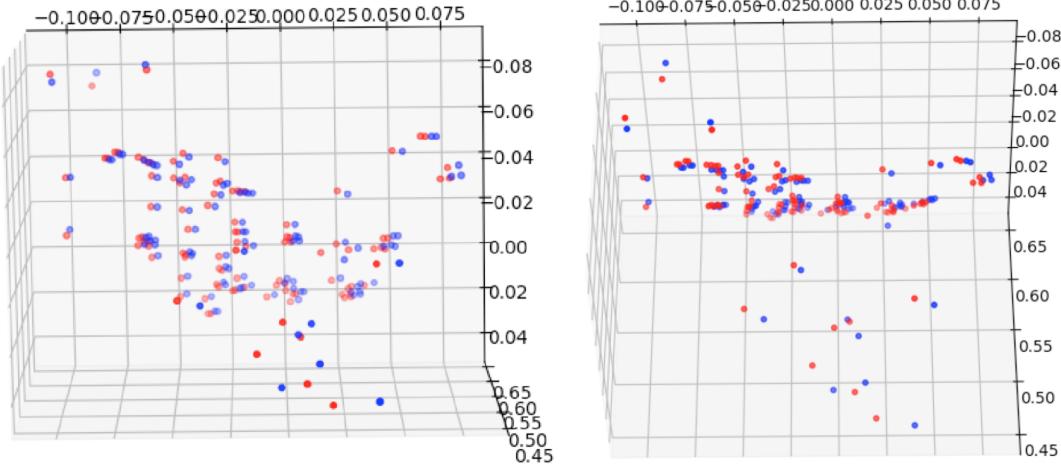


Figure 43: Visualization of 3D points for noisy correspondences before and after using bundle adjustment

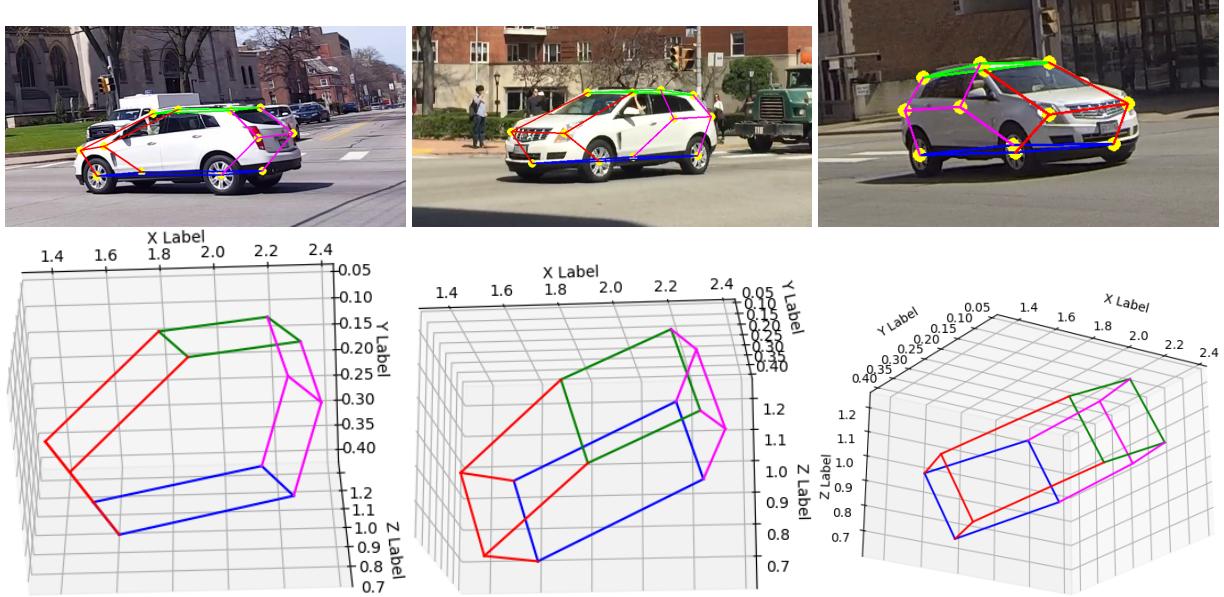


Figure 44: An example detections on the top and the reconstructions from multiple views

6 Multiview Keypoint Reconstruction (Extra Credit)

You will use multi-view capture of moving vehicles and reconstruct the motion of a car. The first part of the problem will be using a single time instance capture from three views (Figure 44 Top) and reconstruct vehicle keypoints and render from multiple views (Figure 44 Bottom). Make use of `q6_ec_multiview_reconstruction.py` file and `data/q6` folder contains the images.

Q6.1 [Extra Credit - 15 points] Write a function to compute the 3D keypoint locations P given the 2D part detections pts1 , pts2 and pts3 and the camera projection matrices $C1$, $C2$, $C3$. The camera matrices are given in the numpy files.

```
[P, err] = MultiviewReconstruction(C1, pts1, C2, pts2, C3, pts3, Thres)
```

The 2D part detections (pts) are computed using a neural network² and correspond to different locations on a car like the wheels, headlights etc. The third column in pts is the confidence of localization of the keypoints. Higher confidence value represents more accurate localization of the keypoint in 2D. To visualize the 2D detections run `visualize_keypoints(image, pts, Thres)` helper function. Thres is defined as the confidence threshold of the 2D detected keypoints. The camera matrices (C) are computed by running an SfM from multiple views and are given in the numpy files with the 2D locations. By varying confidence threshold Thres (i.e. considering only the points above the threshold), we get different reconstruction and accuracy. Try varying the thresholds and analyze its effects on the accuracy of the reconstruction. Save the best reconstruction (the 3D locations of the parts) from these parameters into a `q6_1.npz` file.

Hint: You can modify the triangulation function to take three views as input. After you do the threshold lets say m points lie above the threshold and n points lie below the threshold. Now your task is to use these m good points to compute the reconstruction. For each 3D location use two view or three view triangulation for intialization based on visibility after thresholding.

²Code Used For Detection and Reconstruction

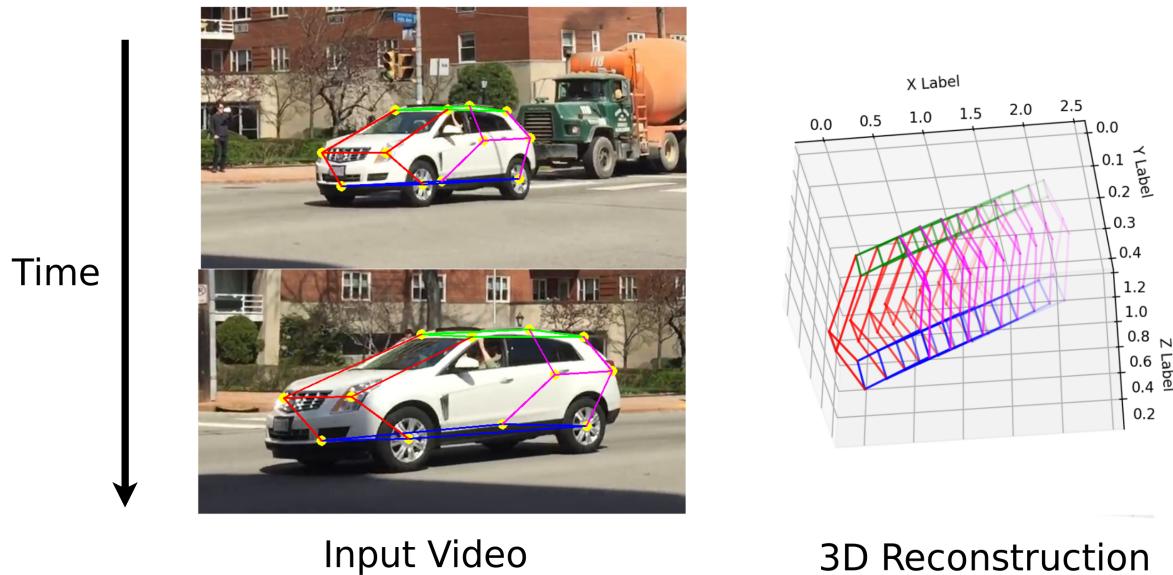


Figure 45: Spatiotemporal reconstruction of the car (right) with the projections at two different time instances in a single view(left)

In your write-up:

- Describe the method you used to compute the 3D locations.
- Include an image of the Reconstructed 3D points with the points connected using the helper function `plot_3d_keypoint (P)` with the reprojection error.
- Include the code snippets `MultiviewReconstruction` in your write-up.

Q6.1

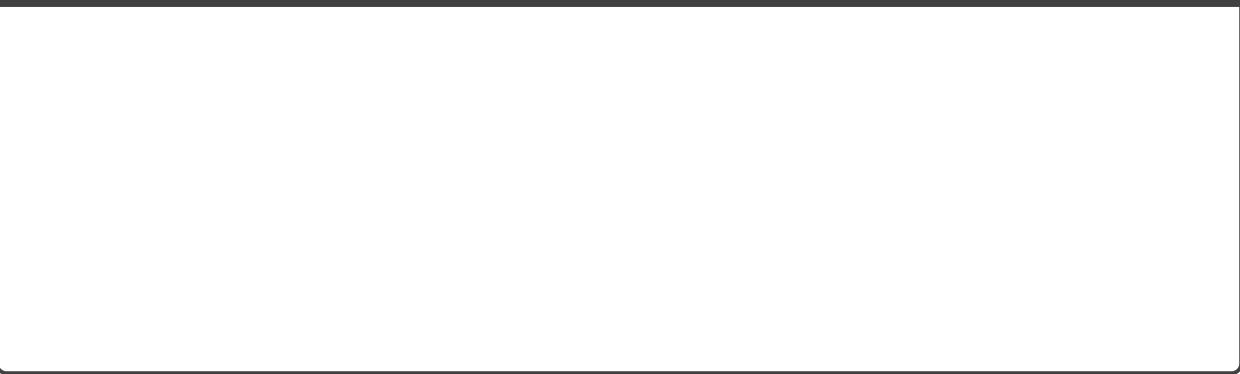
Q6.2 [Extra Credit - 15 points] From the previous question you have done a 3D reconstruction at a time instance. Now you are going to iteratively repeat the process over time and compute a spatio temporal reconstruction of the car. The images in the `data/q6` folder shows the motion of the car at an intersection captured from multiple views. The images are given as (`cam1_time0.jpg`, ..., `cam1_time9.jpg`) for camera 1 and (`cam2_time0.jpg`, ..., `cam2_time9.jpg`) for camera2 and (`cam3_time0.jpg`, ..., `cam3_time9.jpg`) for camera3. The corresponding detections and camera matrices are given in (`time0.npz`, ..., `time9.npz`). Use the above details and compute

the spatio temporal reconstruction of the car for all 10 time instances and plot them by completing the `plot_3d_keypoint_video` function. A sample plot with the first and last time instance reconstruction of the car with the reprojections shown in the Figure 45.

In your write-up:

- Plot the spatio-temporal reconstruction of the car for the 10 timesteps.
- Include the code snippets `plot_3d_keypoint_video` in your write-up.

Q6.2



7 Deliverables

The assignment (code and write-up) should be submitted to Gradescope. The write-up should be named `<AndrewId>.hw3.pdf` and the code should be a zip named `<AndrewId>.hw3.zip`. ***Please make sure that you assign the location of answers to each question on Gradescope.*** The zip should have the following files in the structure defined below. (Note: Neglecting to follow the submission structure will incur a huge score penalty!). You can run the included `checkA4Submission.py` script to ensure that your zip folder structure is correct.

- `<AndrewId>.hw3.pdf`: your write-up.
- `q2_1_eightpoint.py`: script for Q2.1.
- `q2_2_sevenpoint.py`: script for Q2.2.
- `q3_1_essential_matrix.py`: script for Q3.1.
- `q3_2_triangulate.py`: script for Q3.2.
- `q4_1_epipolar_correspondence.py`: script for Q4.1.
- `q4_2_visualize.py`: script for Q4.2.
- `q5_bundle_adjustment.py`: script for Q5.
- `q6_ec_multiview_reconstruction.py`: script for (extra-credit) Q6.
- `helper.py`: helper functions.
- `q2_1.npz`: file with output of Q2.1.
- `q2_2.npz`: file with output of Q2.2.

- q3_1.npz: file with output of Q3.1.
- q3_3.npz: file with output of Q3.3.
- q4_1.npz: file with output of Q4.1.
- q4_2.npz: file with output of Q4.2.
- q6_1.npz: (extra-credit) file with the output of Q6.1.

***Do not include the data directory in your submission.**

8 FAQs

Credits: Paul Nadan

Q2.1: Does it matter if we unscale \mathbf{F} before or after calling refineF?

The relationship between \mathbf{F} and $\mathbf{F}_{normalized}$ is fixed and defined by a set of transformations, so we can convert at any stage before or after refinement. The nonlinear optimization in refineF may work slightly better with normalized \mathbf{F} , but it should be fine either way.

Q2.1: Why does the other image disappear (or become really small) when I select a point using the displayEpipolarF GUI?

This issue occurs when the corresponding epipolar line to the point you selected lies far away from the image. Something is likely wrong with your fundamental matrix.

Q2.1 Note: The GUI will provide the correct epipolar lines even if the program is using the wrong order of pts1 and pts2 in calculating the eightpoint algorithm. So one thing to check is that the optimizer should only take < 10 iterations (shown in the output) to converge if the ordering is correct.

Q3.2: How can I get started formulating the triangulation equations?

One possible method: from the first camera, $x_{1i} = P_1\omega_1 \implies x_{1i} \times P_1\omega_1 = 0 \implies A_{1i}\omega_i = 0$. This is a linear system of 3 equations, one of which is redundant (a linear combination of the other two), and 4 variables. We get a similar equation from the second camera, for a total of 4 (non-redundant) equations and 4 variables, i.e. $A_i\omega_i = 0$.

Q3.2: What is the expected value of the reprojection error?

The reprojection error for the data in `some_corresp.npz` should be around 352 (or 89 without using refineF). If you get a reprojection error of around 94 (or 1927 without using refineF) then you have somehow ended up with a transposed \mathbf{F} matrix in your eightpoint function.

Q3.2: If you are getting high reprojection error but can't find any errors in your triangulate function?

one useful trick is to temporarily comment out the call to refineF in your 8-point algorithm and make sure that the epipolar lines still match up. The refineF function can sometimes find a pretty good solution even starting from a totally incorrect matrix, which results in the \mathbf{F} matrix passing the sanity checks even if there's an error in the 8-point function. However, having a slightly incorrect \mathbf{F} matrix can still cause the reprojection error to be really high later on even if your triangulate code is correct.

Q4.2 Note: Figure 7 in the assignment document is incorrect - if you look closely you'll notice that the z coordinates are all negative. Don't worry if your solution is different from the example as long as the 3D structure of the temple is evident.

Q5.1: How many inliers should I be getting from RANSAC?

The correct number of inliers should be around 106. This provides a good sanity check for whether the chosen tolerance value is appropriate.

References

- [1] David A Forsyth and Jean Ponce. *Computer vision: a modern approach*. prentice hall professional technical reference, 2002.
- [2] Richard Szeliski. *Computer vision: algorithms and applications*. Springer Nature, 2022.